

UMAP AND T-SNE ARE ALSO TWO USEFUL TECHNIQUES TO VISUALIZE LARGE DIMENSIONAL SPACES

INSTEAD OF MATRIX FACTORIZATION SUCH AS PCA, THEY
ARE BASED ON GRAPH LAYOUT

1. WE BUILD A GRAPH OF OUR HIGH-DIMENSIONALITY DATA
2. WE OPTIMIZE TO GET THE MOST SIMILAR GRAPH IN TWO DIMENSIONS

[https://
lvdmaaten.github.io/tsne/](https://lvdmaaten.github.io/tsne/)

https://umap-learn.readthedocs.io/en/latest/how_umap_works.html

UMAP EXAMPLES AND EXPLANATIONS;

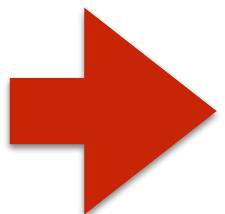
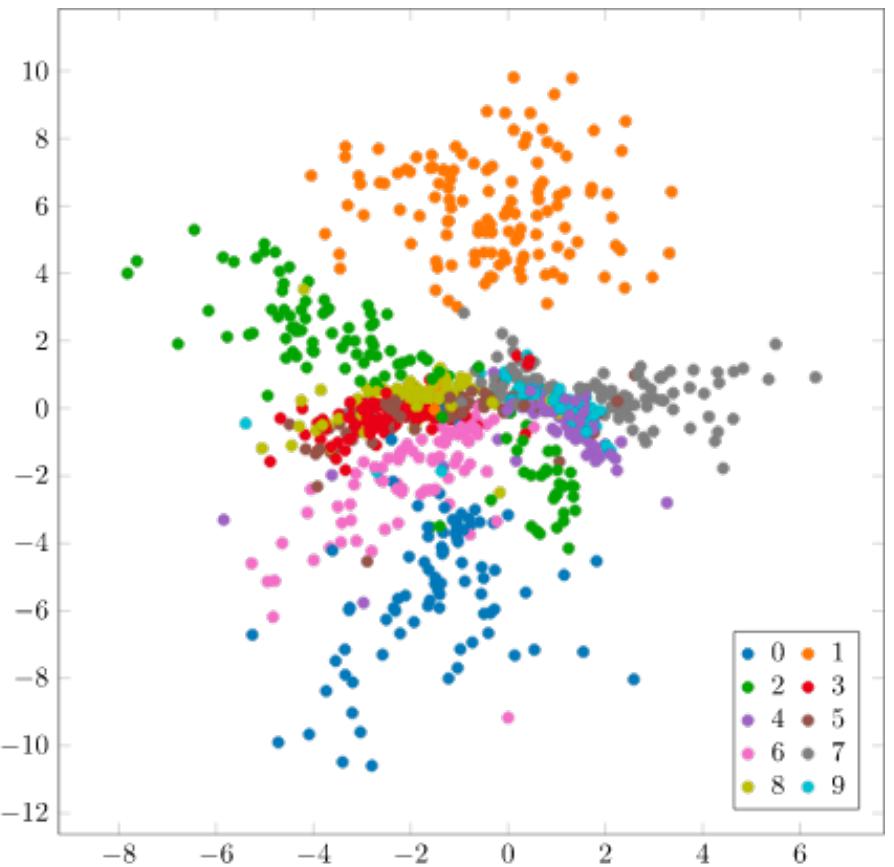
<https://pair-code.github.io/understanding-umap/>

[#:~:text=In%20the%20simplest%20sense%2C%20UMAP,behind%20them%](#)
[20is%20remarkably%20simple.](#)

HOW MIGHT YOU SOLVE THESE
RELATED PROBLEMS ONCE YOU HAVE
THE LATENT SPACE COORDINATES FOR
YOUR TRAINING SAMPLE?

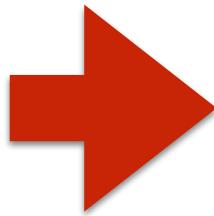
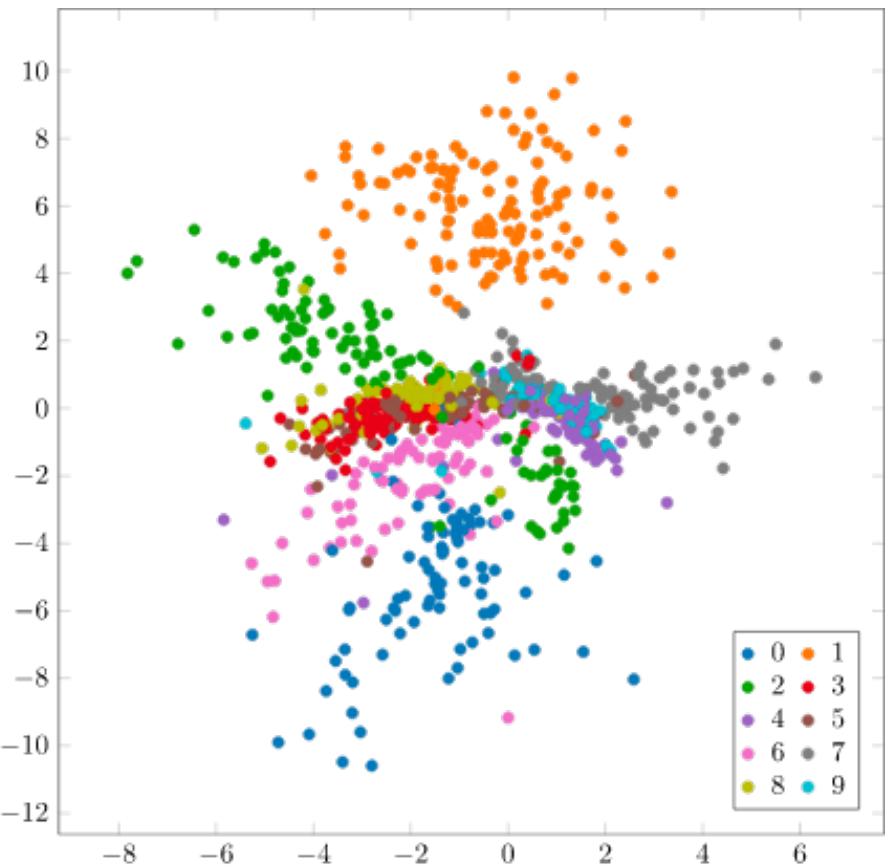
GENERATE A RANDOM SAMPLE DRAWN FROM THE INPUT
DISTRIBUTION ("**GENERATIVE MODEL**")

ESTIMATE THE PROBABILITY DENSITY OF AN ARBITRARY
INPUT, RELATIVE TO THE INPUT DISTRIBUTION
 ("**PROBABILISTIC MODEL**")



HOW CAN I
ESTIMATE $P(X)$?

$(P(z|x))$

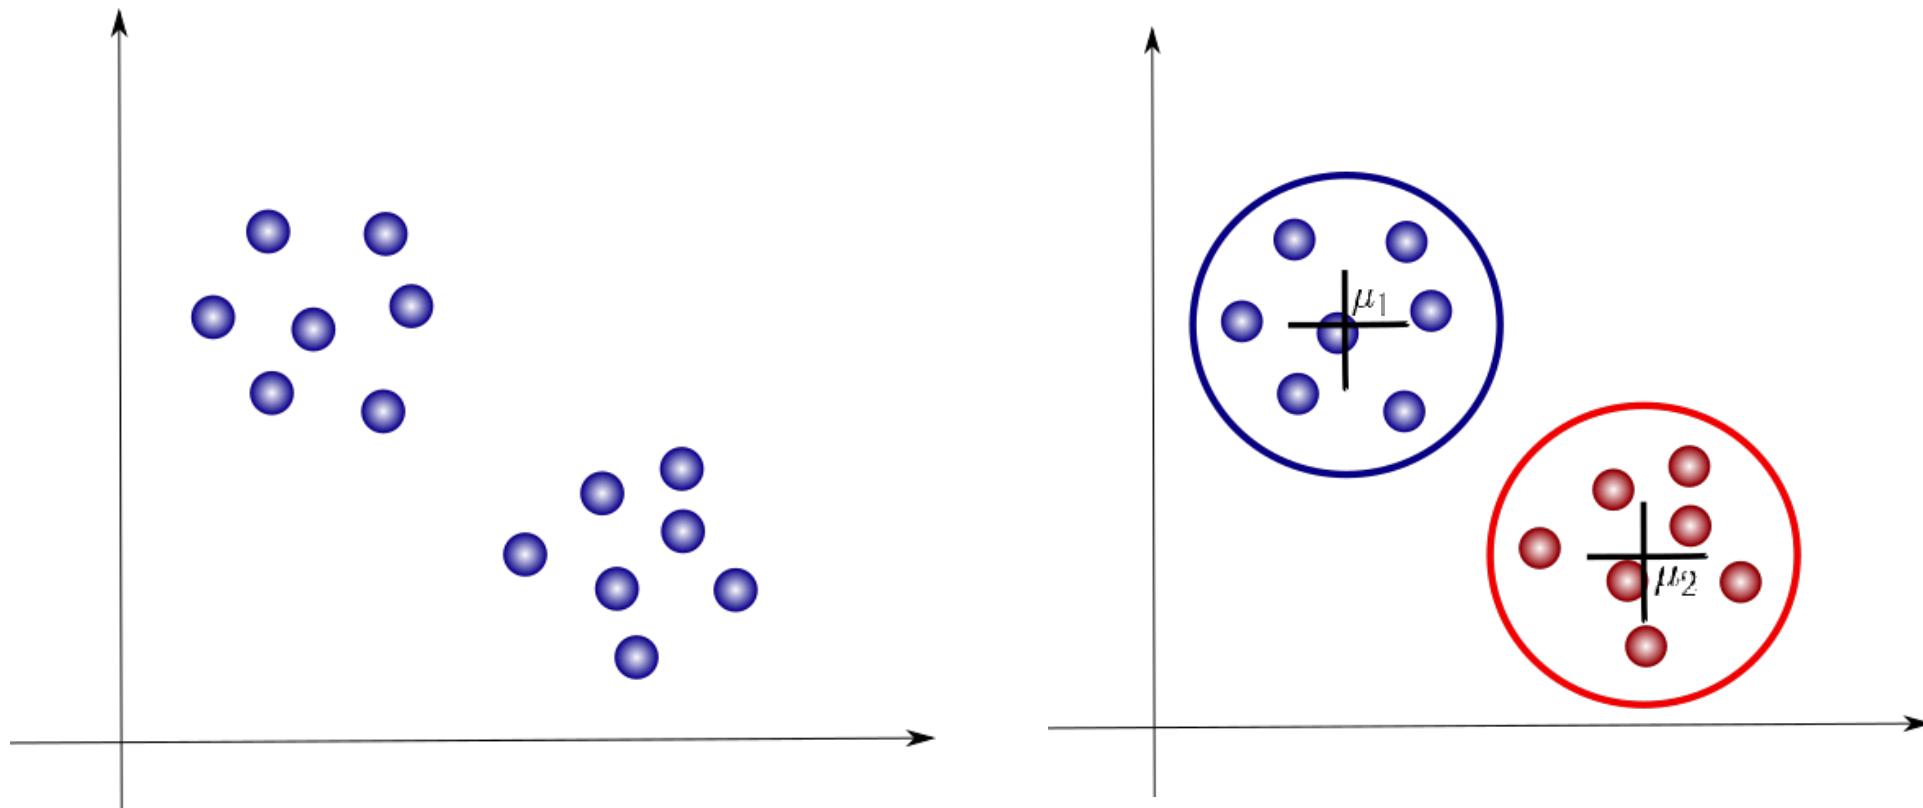


HOW CAN I
ESTIMATE $P(X)$?

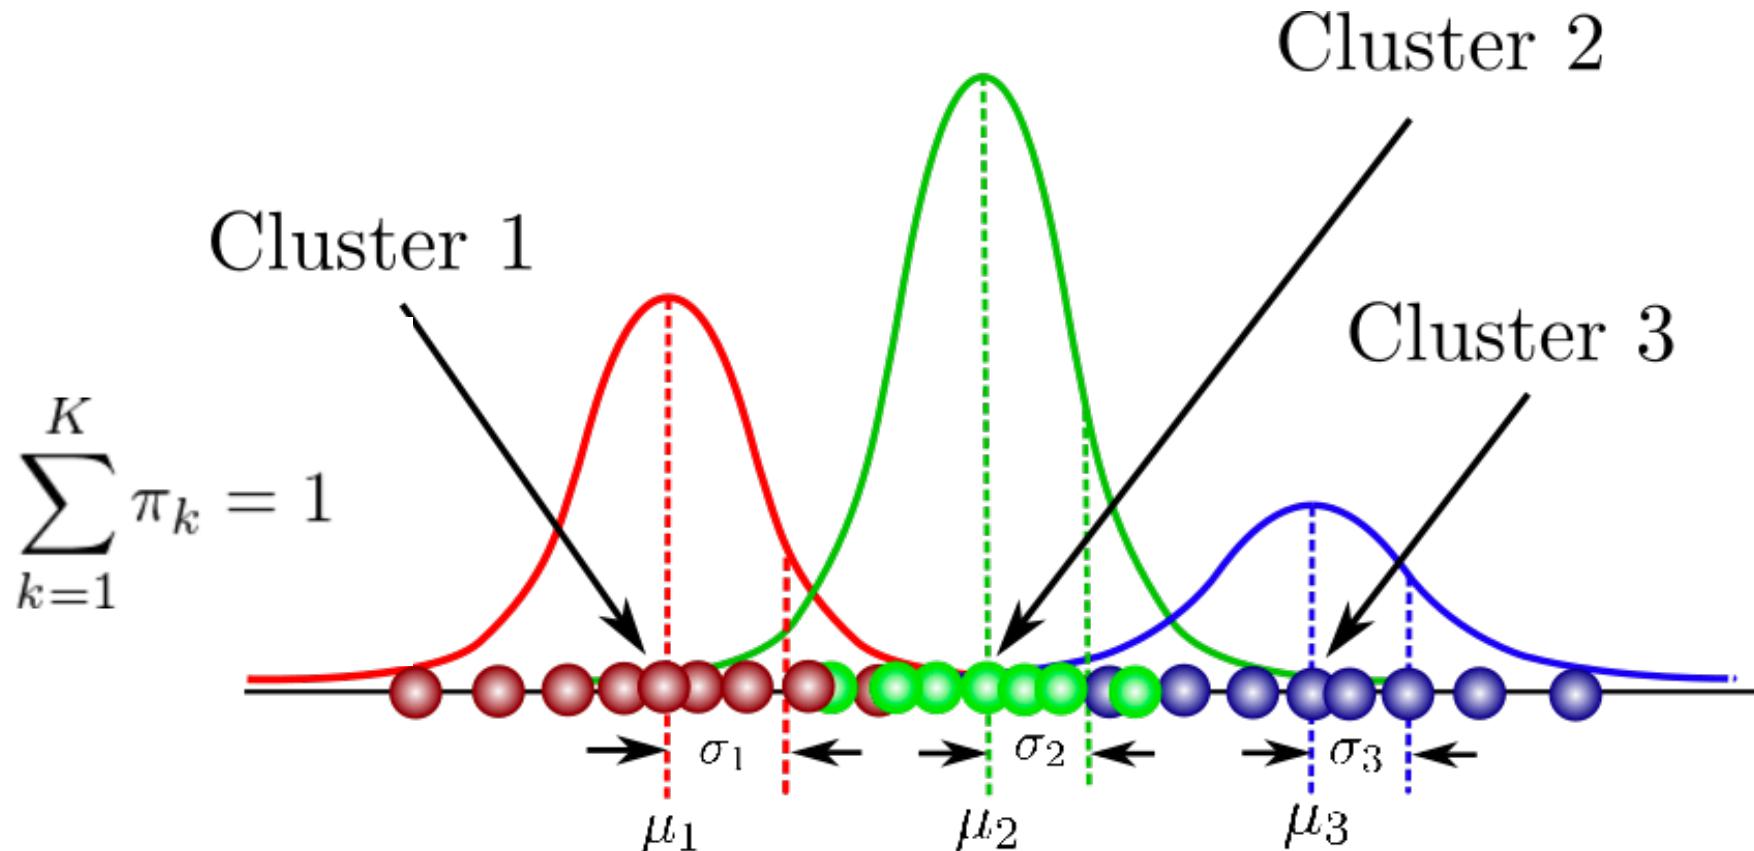
$(P(z|x))$

WHEN YOU DON'T KNOW, ASSUME IT IS GAUSSIAN...

GAUSSIAN MIXTURE MODELS (GMMs) ARE DENSITY ESTIMATOR METHODS THAT FIT MULTIPLE GAUSSIANS TO THE REPRESENTATION

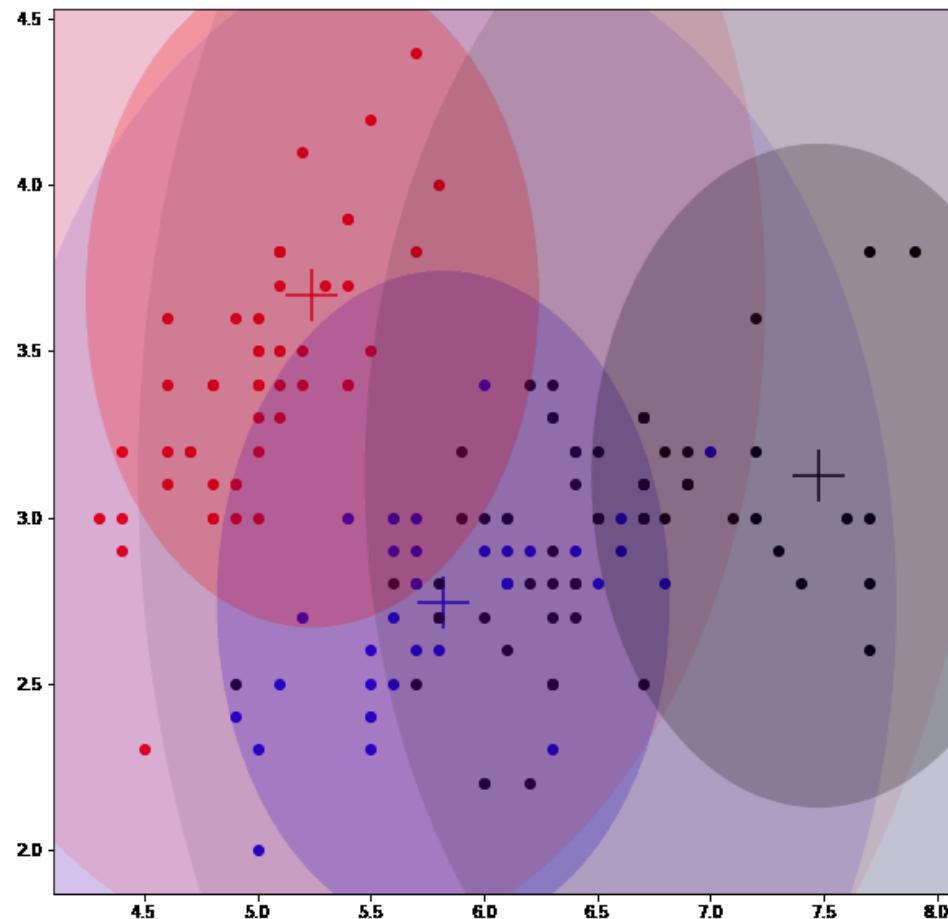


GAUSSIAN MIXTURE MODELS (GMMs) ARE DENSITY ESTIMATOR METHODS THAT FIT MULTIPLE GAUSSIANS TO THE REPRESENTATION

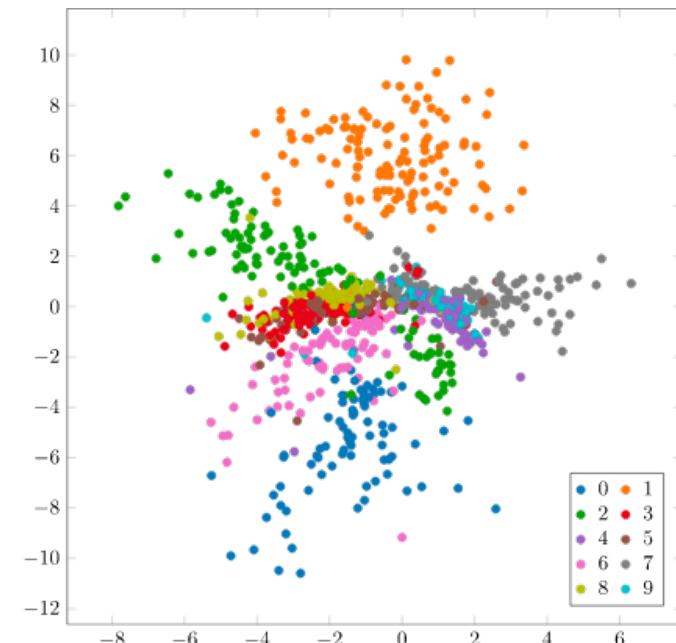
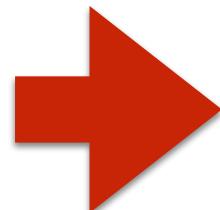
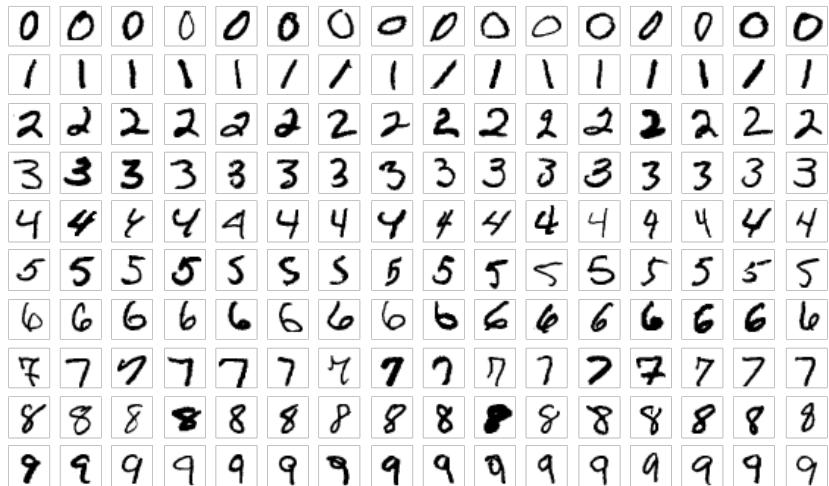


means, sigmas and scale factors of each gaussian are free parameters

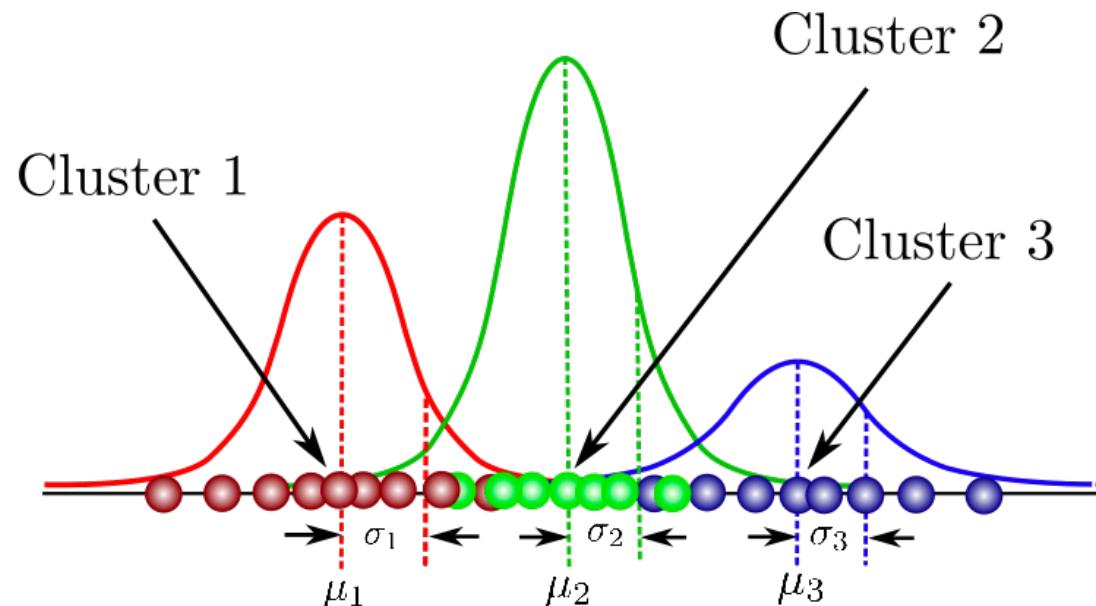
GAUSSIAN MIXTURE MODELS (GMMs) ARE DENSITY ESTIMATOR METHODS THAT FIT MULTIPLE GAUSSIANS TO THE REPRESENTATION



DATA



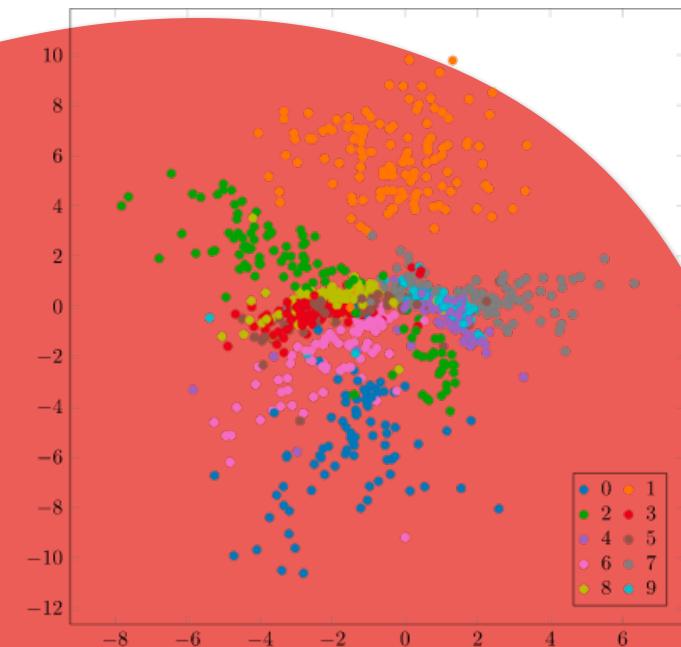
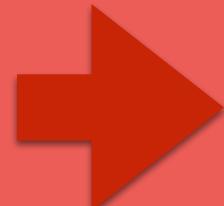
REPRESENTATION
(AUTOENCODER)



MODELING OF $P(X)$ WITH GMMs

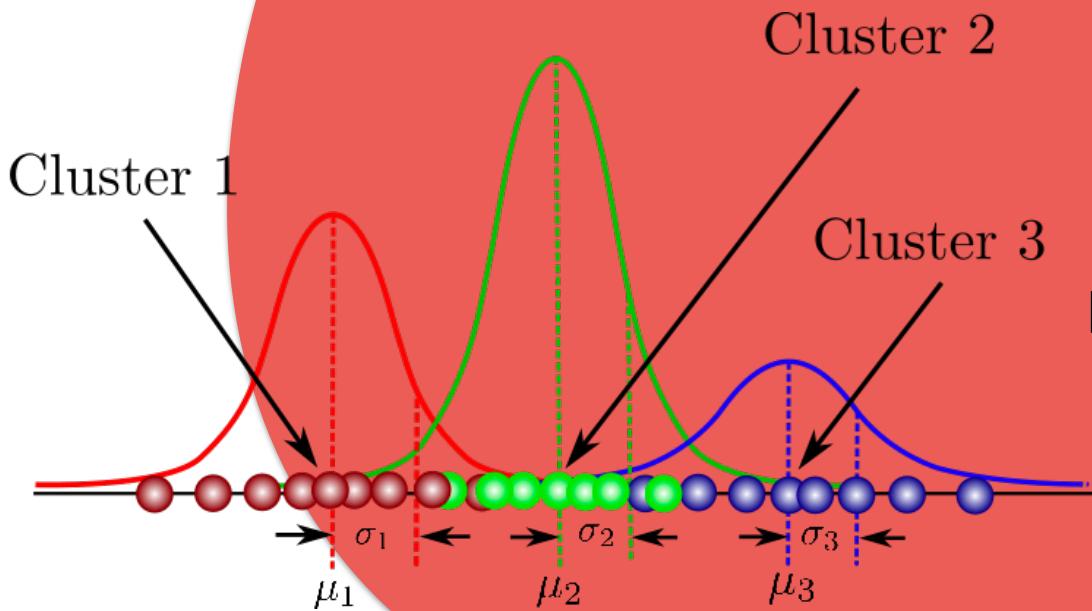
DATA

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3
4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4
5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5
6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6
7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7
8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8
9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9



WE COMBINE THESE 2 STEPS?

REPRESENTATION
(AUTOENCODER)



MODELING OF $P(X)$ WITH GMMs

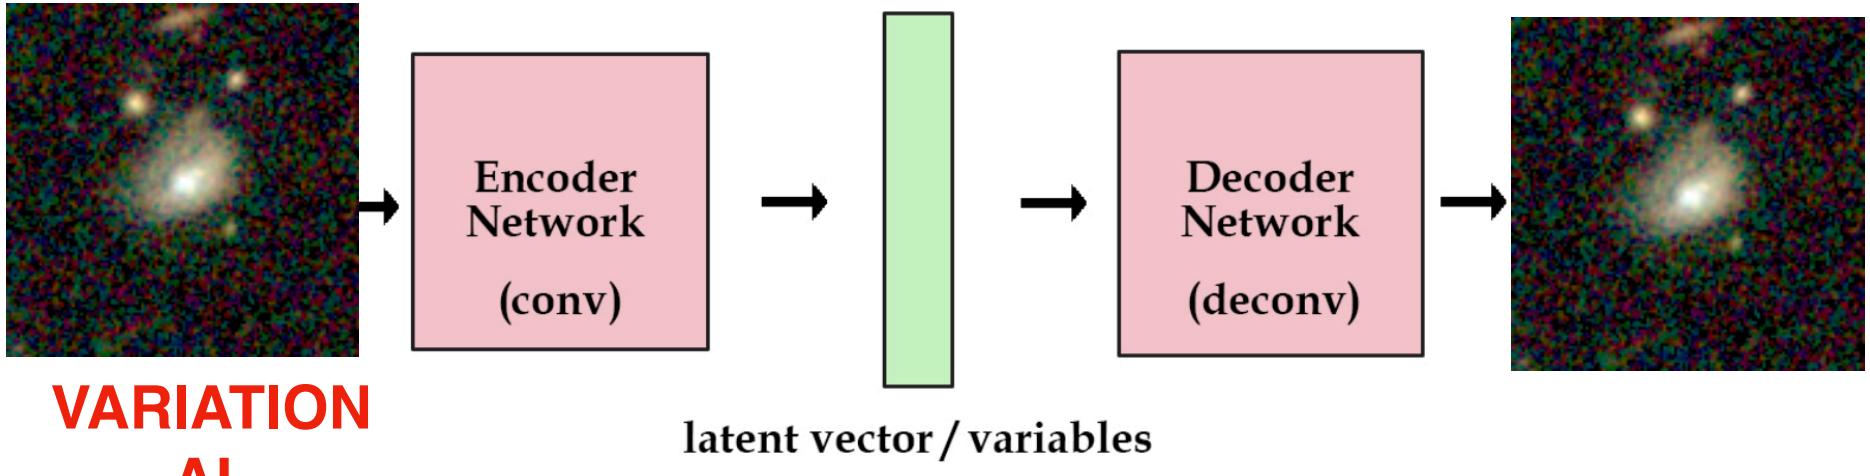
Much of the recent progress in unsupervised deep learning has been to invent network architectures that are capable of solving either or both of these related problems directly, without resorting to any auxiliary methods

VAE
(VARIATIONAL AUTOENCODER)

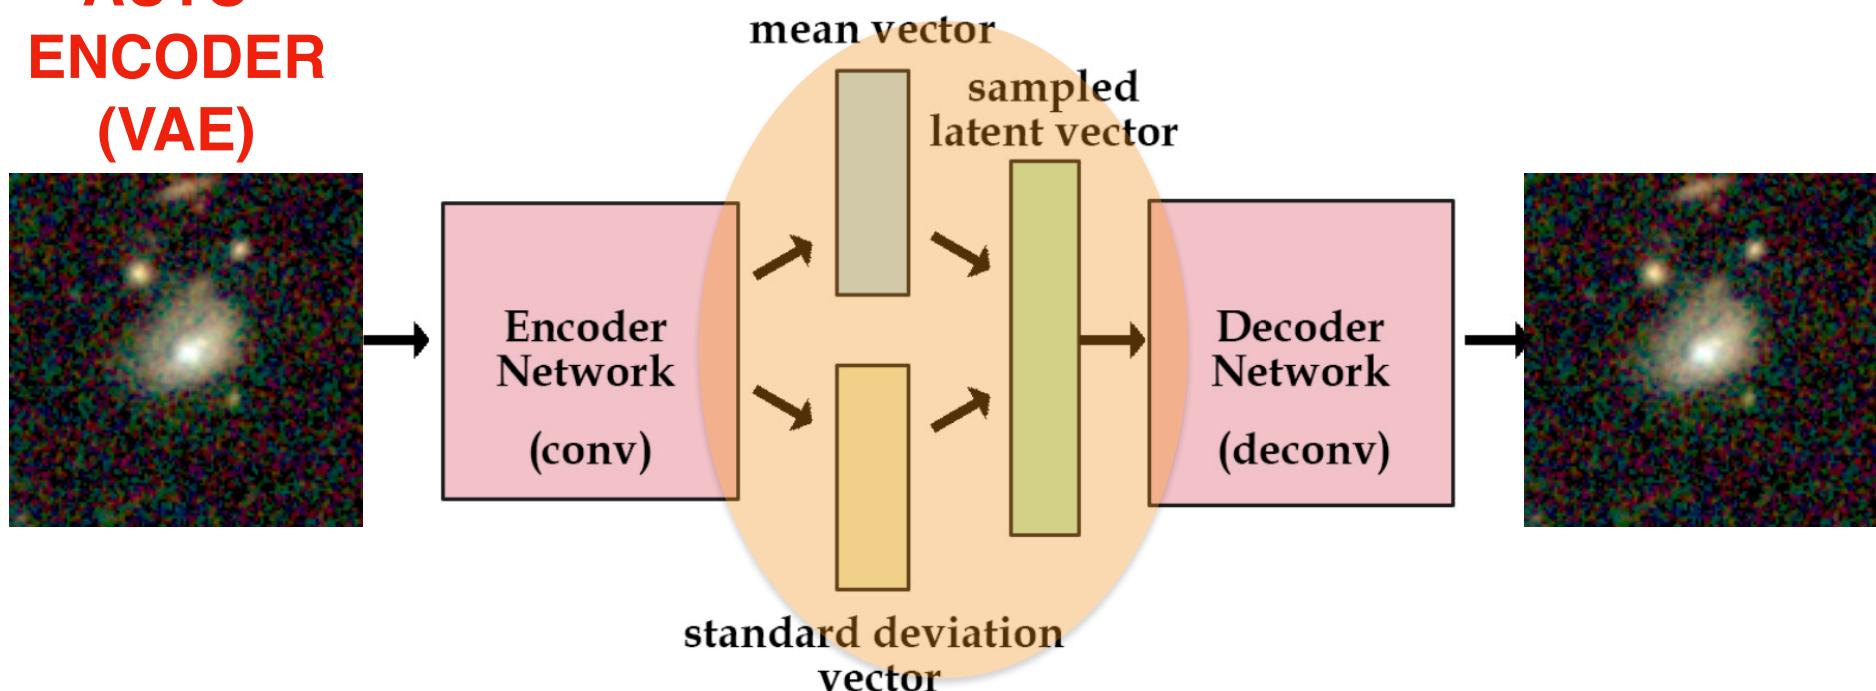
GAN
(GENERATIVE ADVERSARIAL NETWRK)

NF-ARF
(NORMALIZING FLOWS, AUTOREGRESSIVE FLOWS)

AUTO-ENCODER

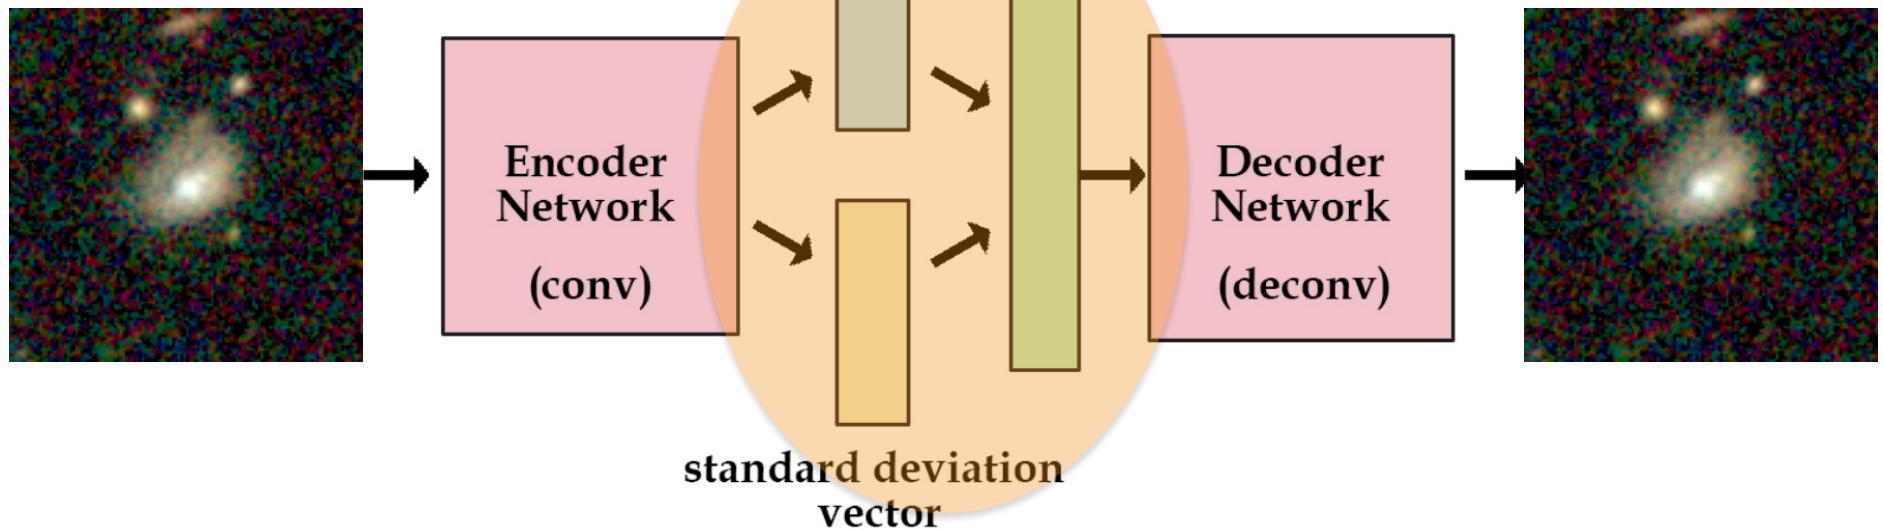


VARIATION
AL
AUTO-
ENCODER
(VAE)



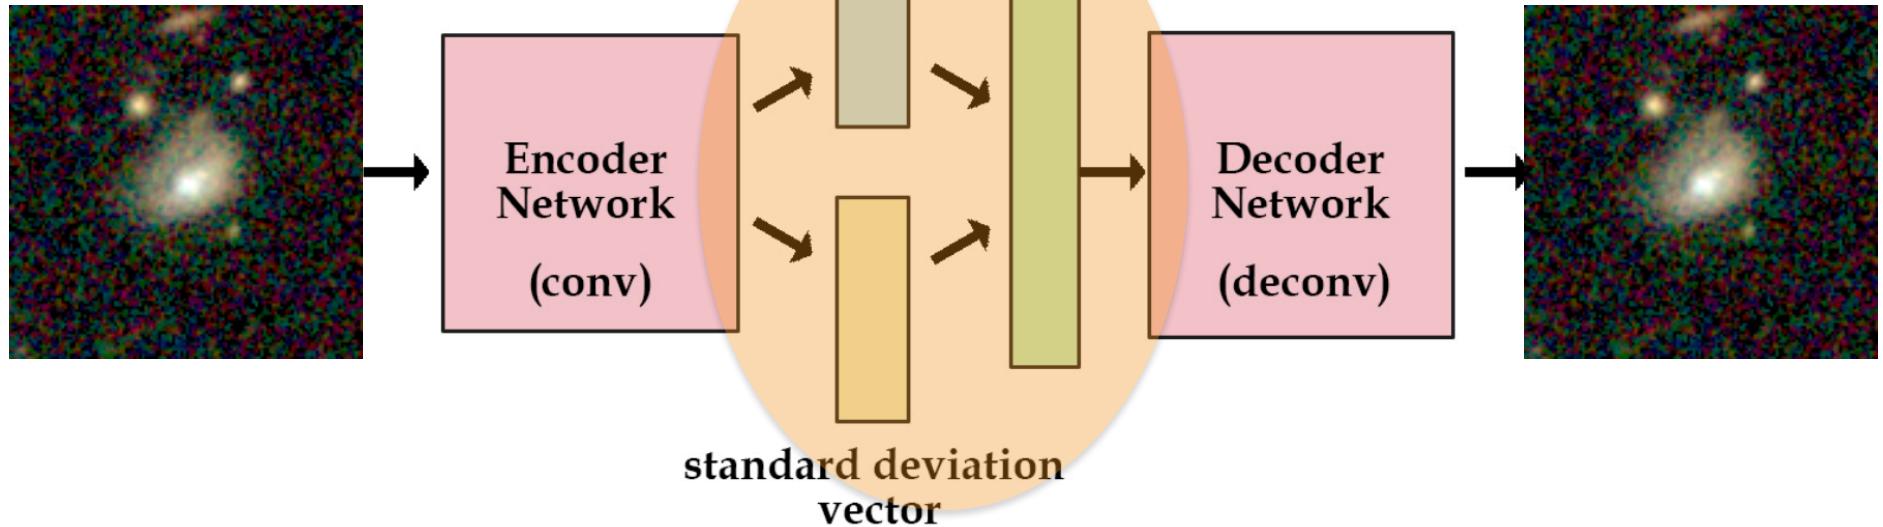
LET'S MODEL THE LATENT SPACE WITH A MIXTURE OF GAUSSIANS

VARIATIONAL AUTO- ENCODER (VAE)



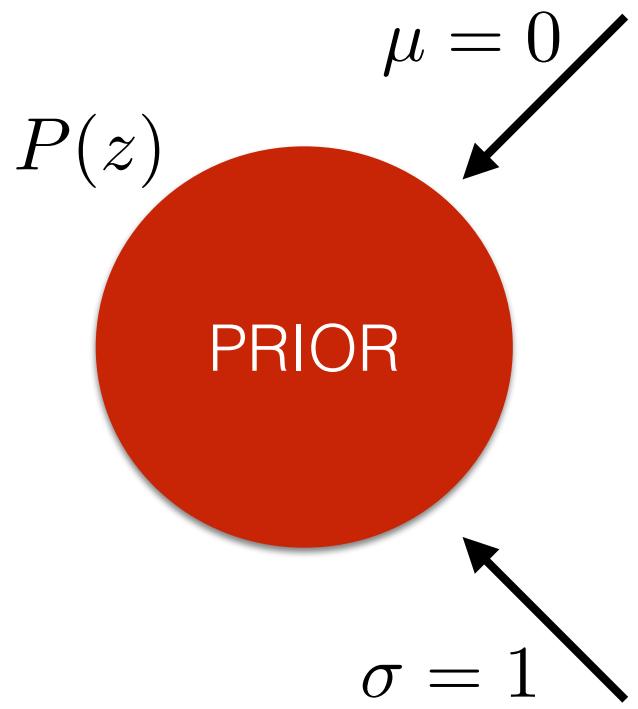
HOWEVER, NOTHING GUARANTEES US THAT THE LATENT SPACE CAN BE MODELLED BY A MIXTURE OF GAUSSIANS....

VARIATIONAL AUTO- ENCODER (VAE)

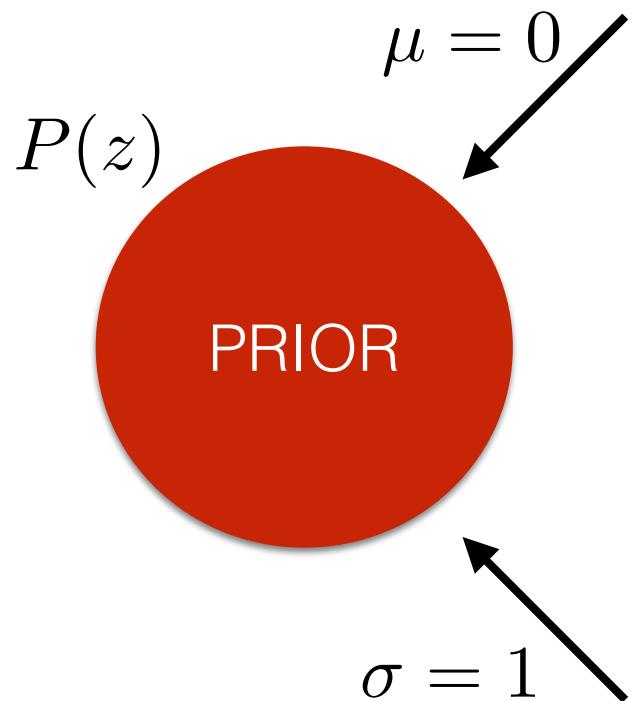


HOWEVER, NOTHING GUARANTEES US THAT THE LATENT SPACE CAN BE MODELLED BY A MIXTURE OF GAUSSIANS....

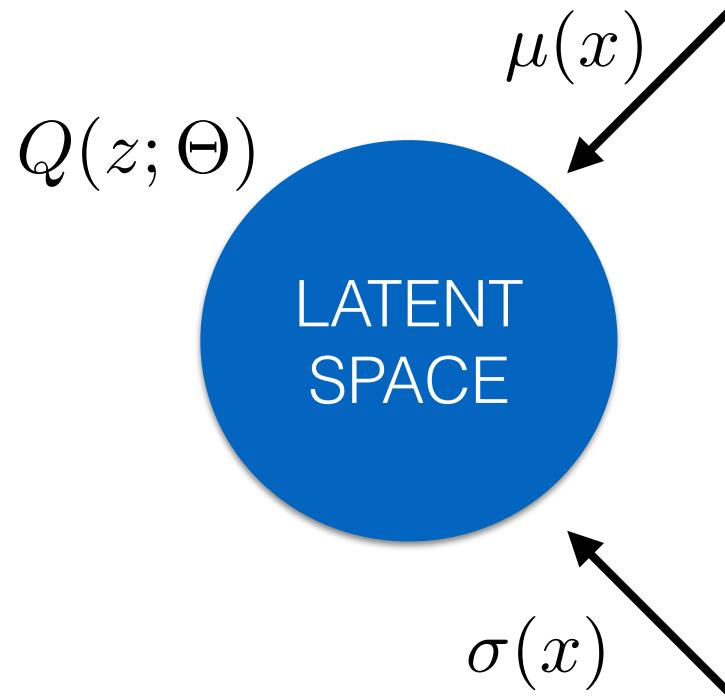
... LET'S FORCE IT TO BE GAUSSIAN LIKE!



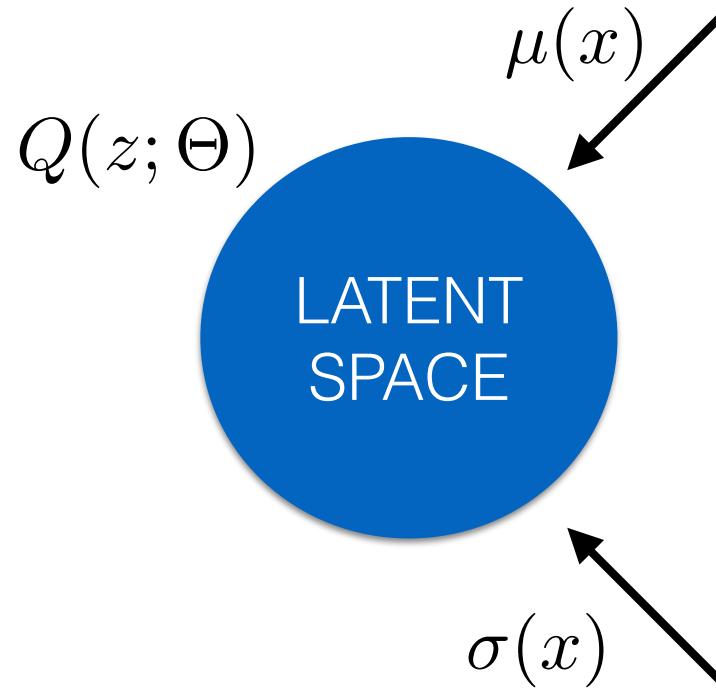
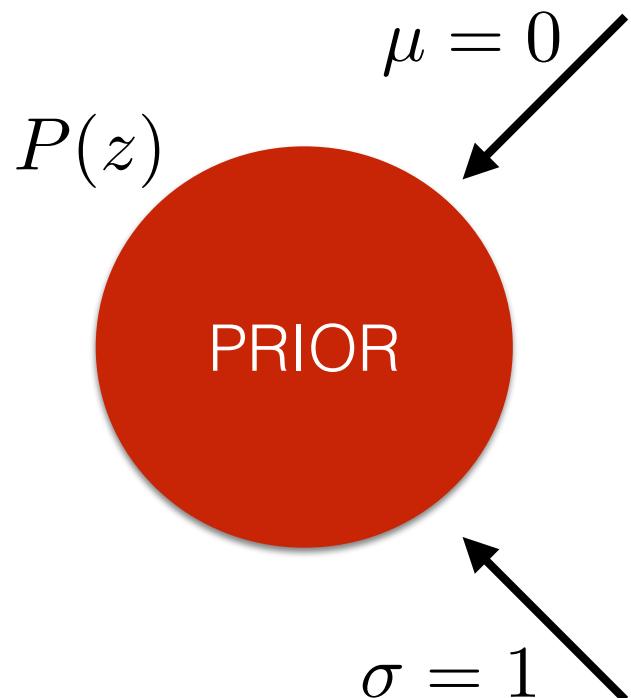
WE ASSUME A SIMPLE PRIOR



WE ASSUME A SIMPLE PRIOR



LATENT SPACE MODELING

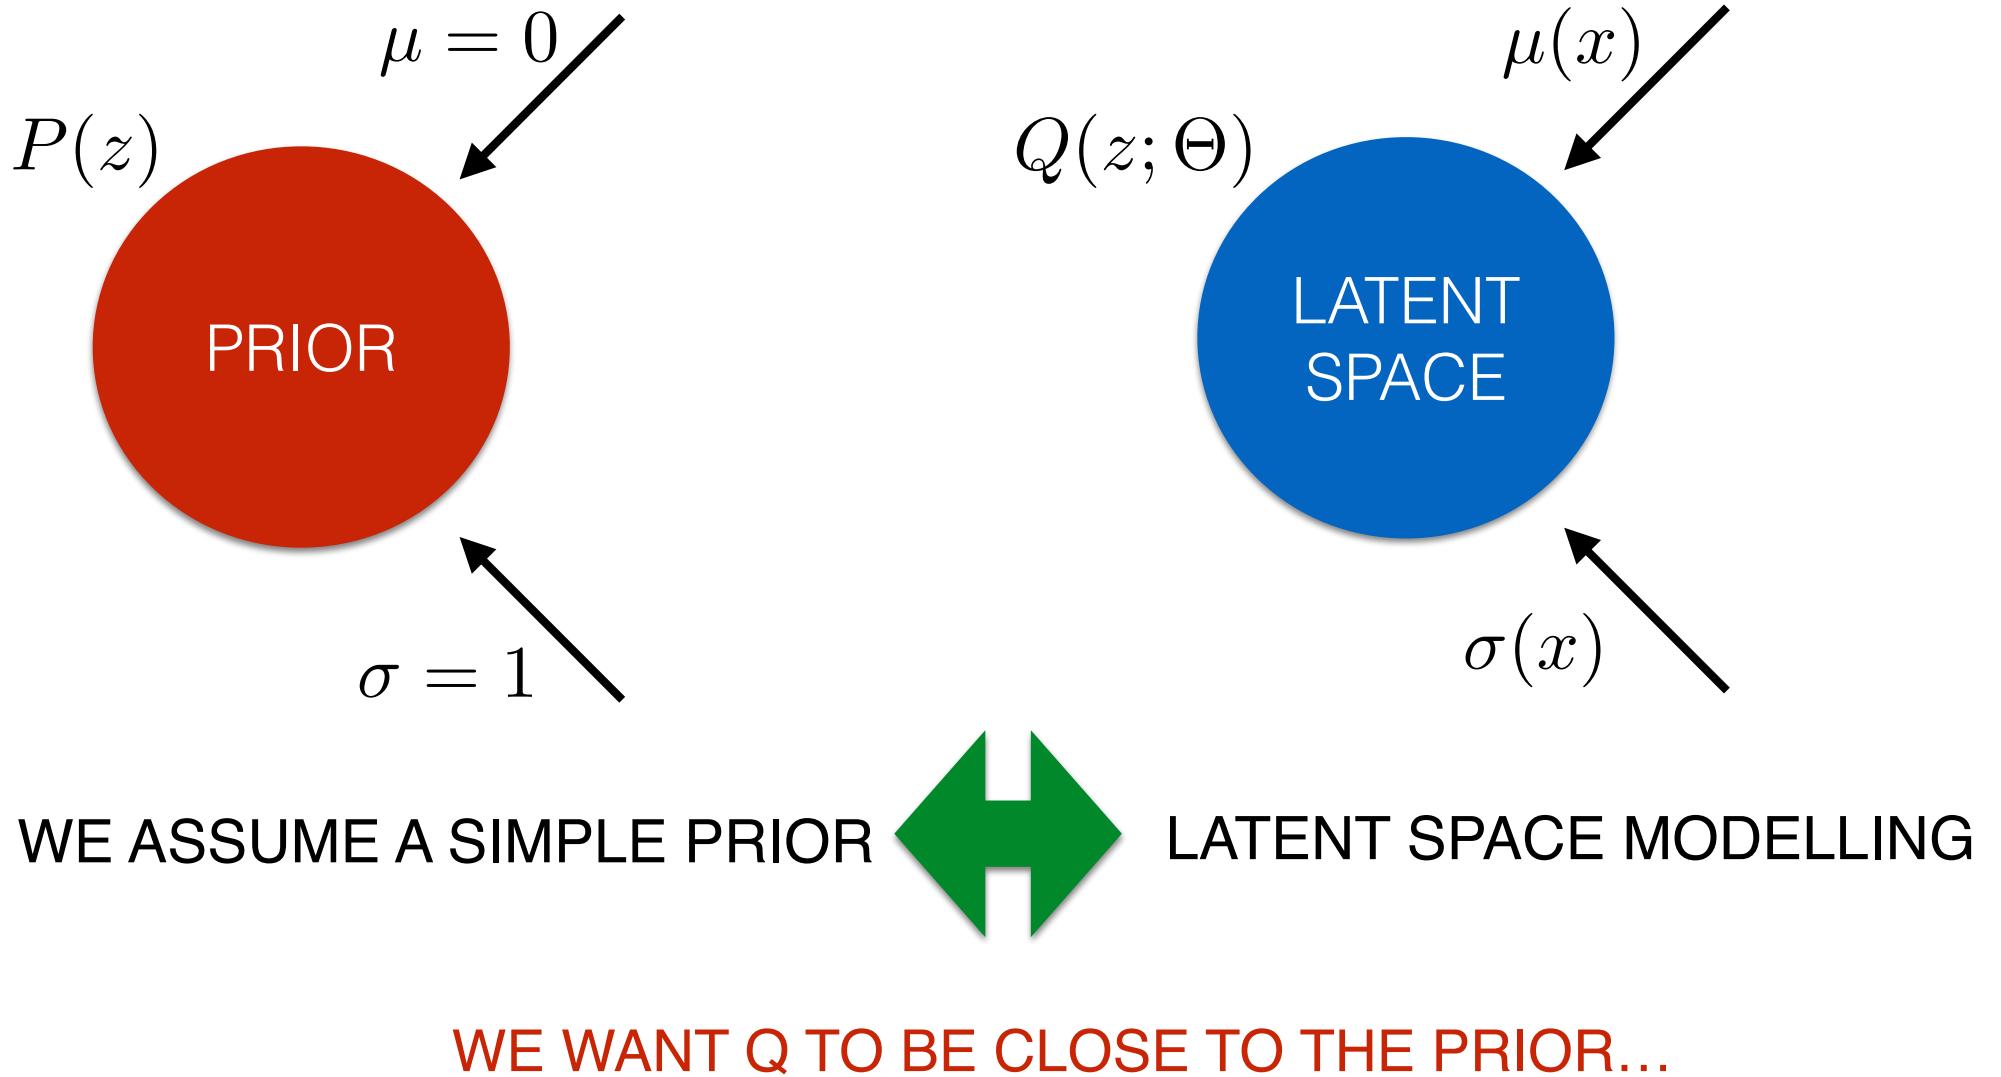


WE ASSUME A SIMPLE PRIOR



LATENT SPACE MODELLING

WE WANT Q TO BE CLOSE TO THE PRIOR...

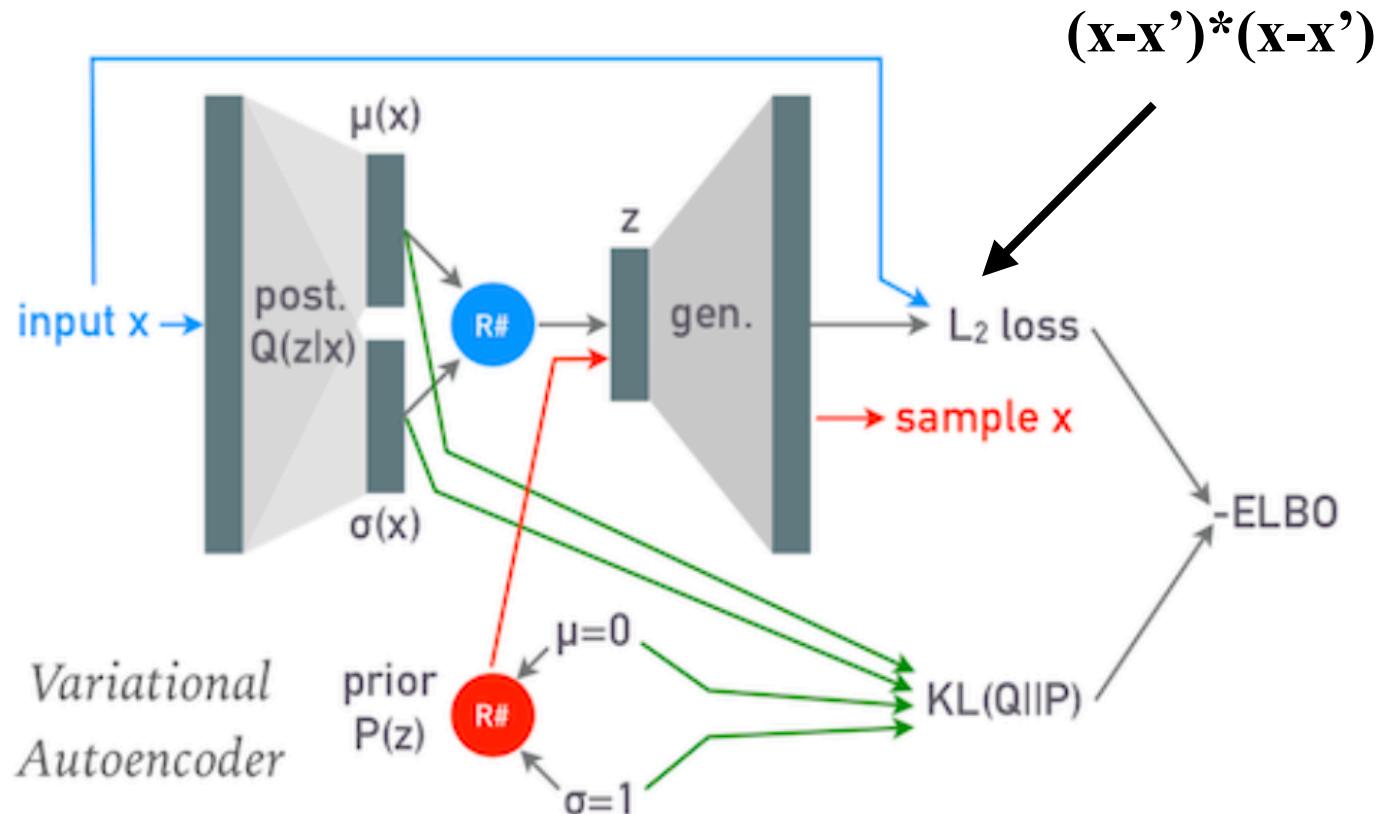


$$D_{\text{KL}}(P \parallel Q) = \sum_{x \in \mathcal{X}} P(x) \log \left(\frac{P(x)}{Q(x)} \right).$$

$$D_{\text{KL}}(P \parallel Q) = \int_{\mathcal{X}} \log \left(\frac{dP}{dQ} \right) \frac{dP}{dQ} dQ,$$

WE MINIMIZE THE K-L DIVERGENCE BETWEEN P AND Q

WHAT WOULD BE THEN THE LOSS FUNCTION OF A VAE?



The key insight of VAE is that we are actually performing variational inference here, which then tells us what the loss function should be...

$$-\text{ELBO} = \langle \log P(\mathbf{x} | \mathbf{z}) \rangle_{\mathbf{z} \sim Q} + \text{KL}(Q(\mathbf{z}; \Theta) \parallel P(\mathbf{z})) ,$$

L2 LOSS

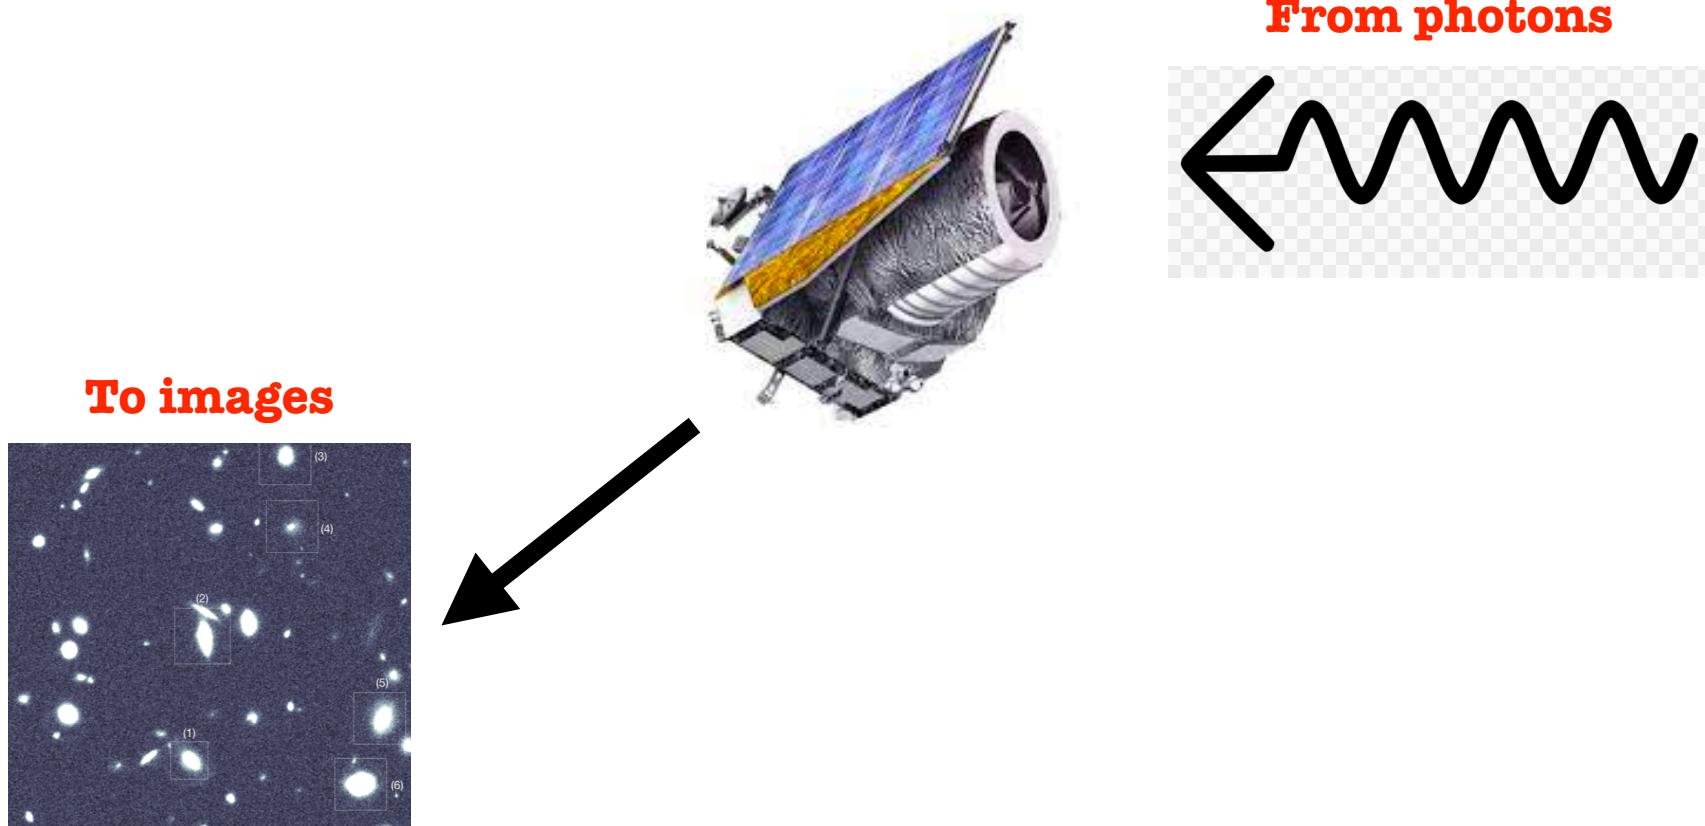
REGULARIZATION TERM

(VQ-VAE)



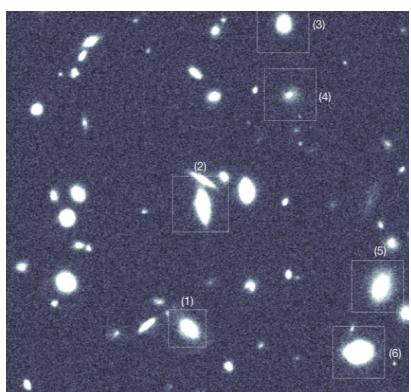
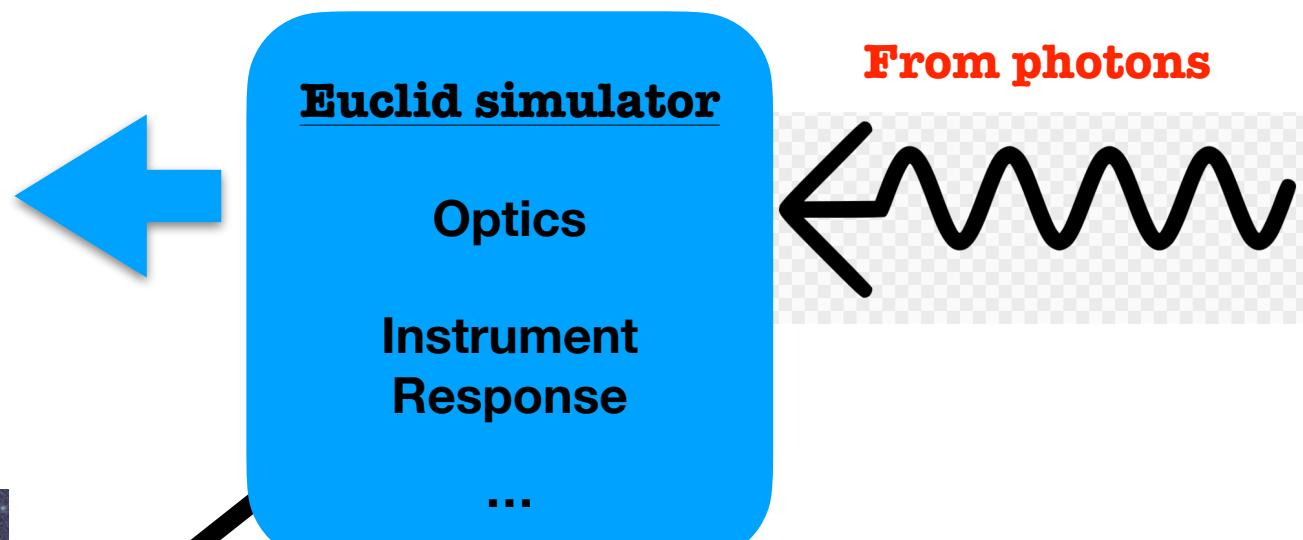
Razavi+19
(deepmind)

Preparing a space mission requires very detailed simulations



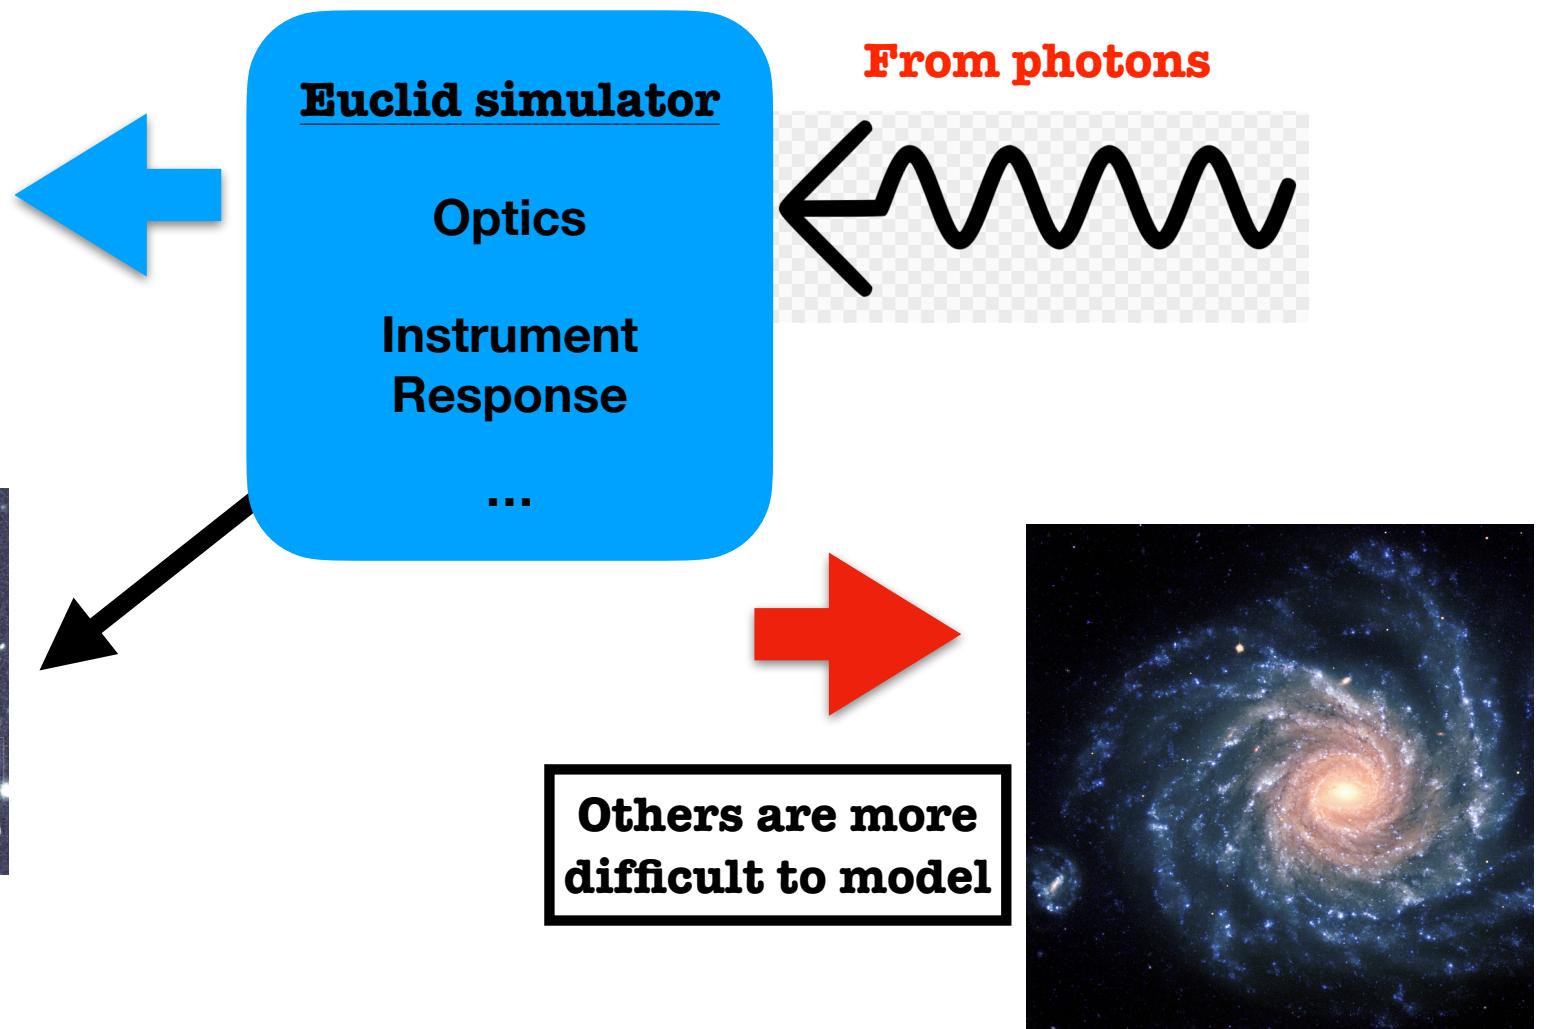
Preparing a space mission requires very detailed simulations

There is a physical model for many of these effects



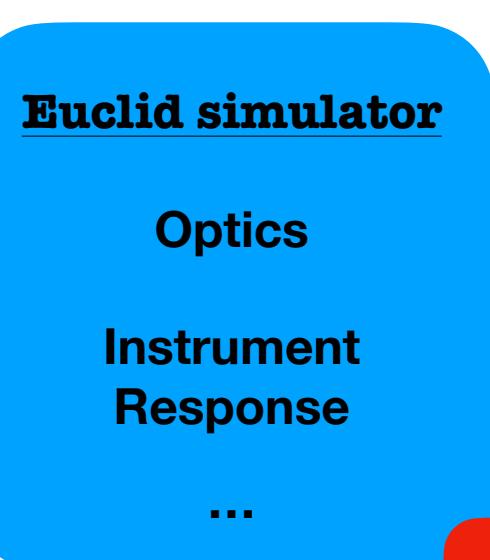
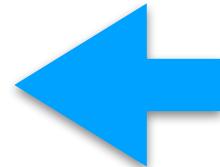
Preparing a space mission requires very detailed simulations

There is a physical model for many of these effects

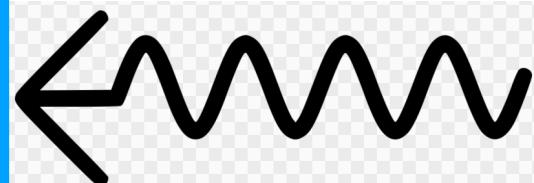


Preparing a space mission requires very detailed simulations

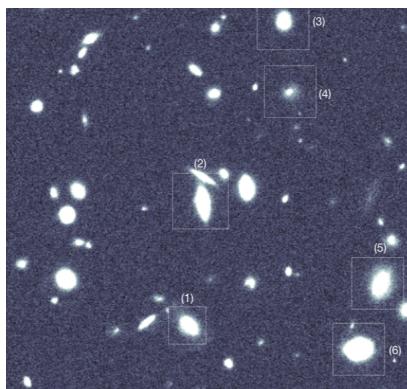
There is a physical model for many of these effects



From photons



To images

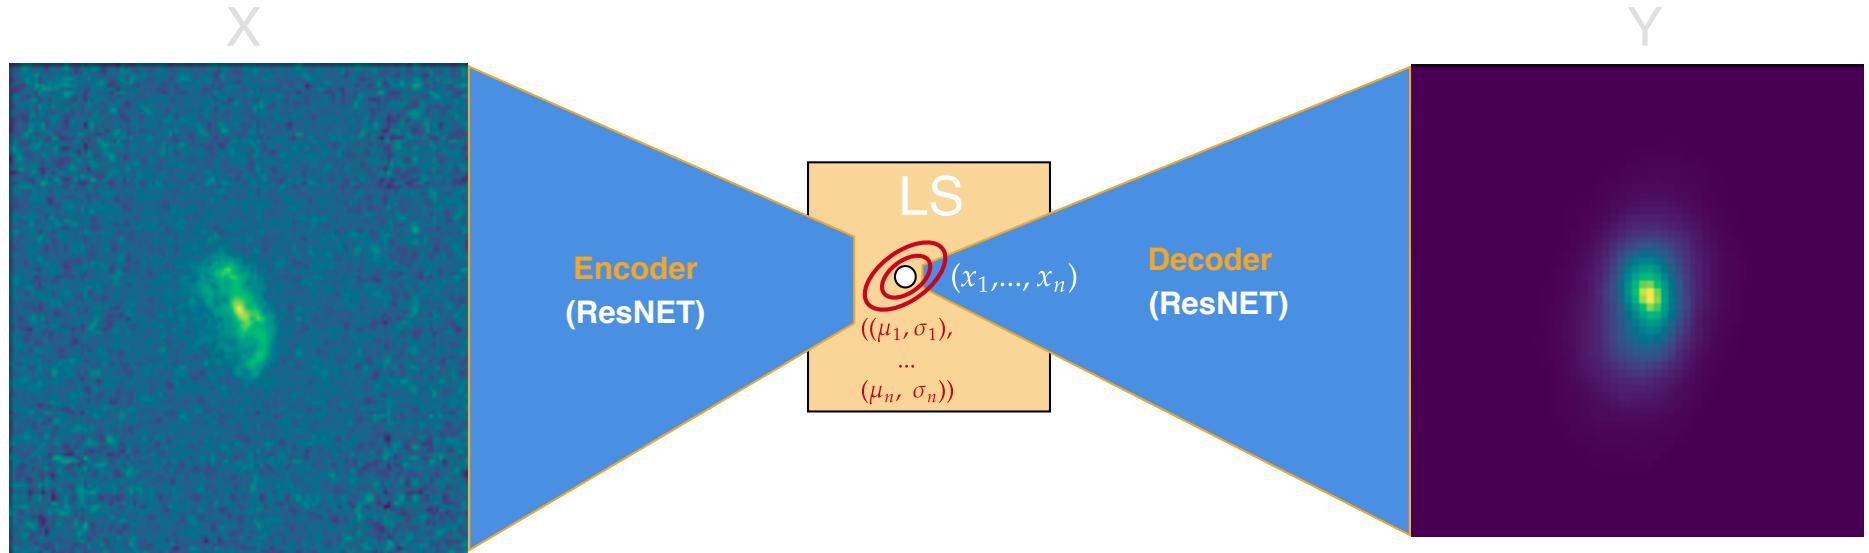


Neural Network

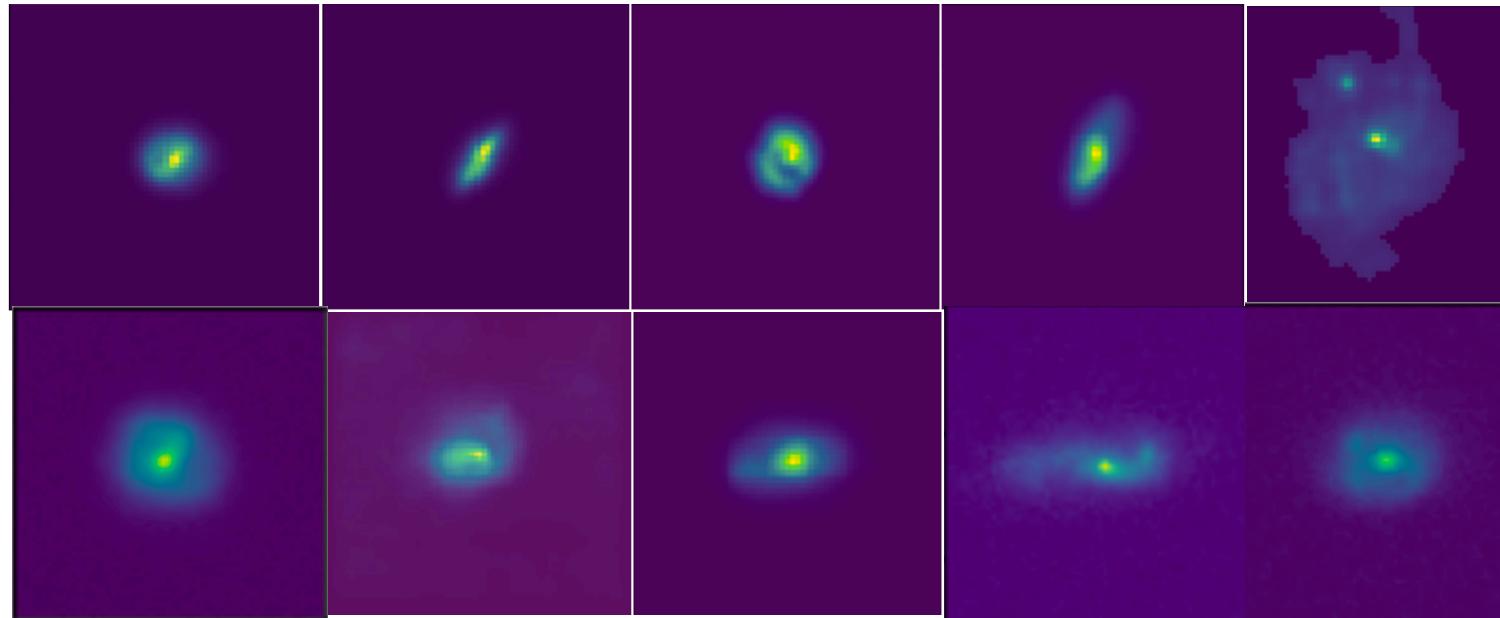
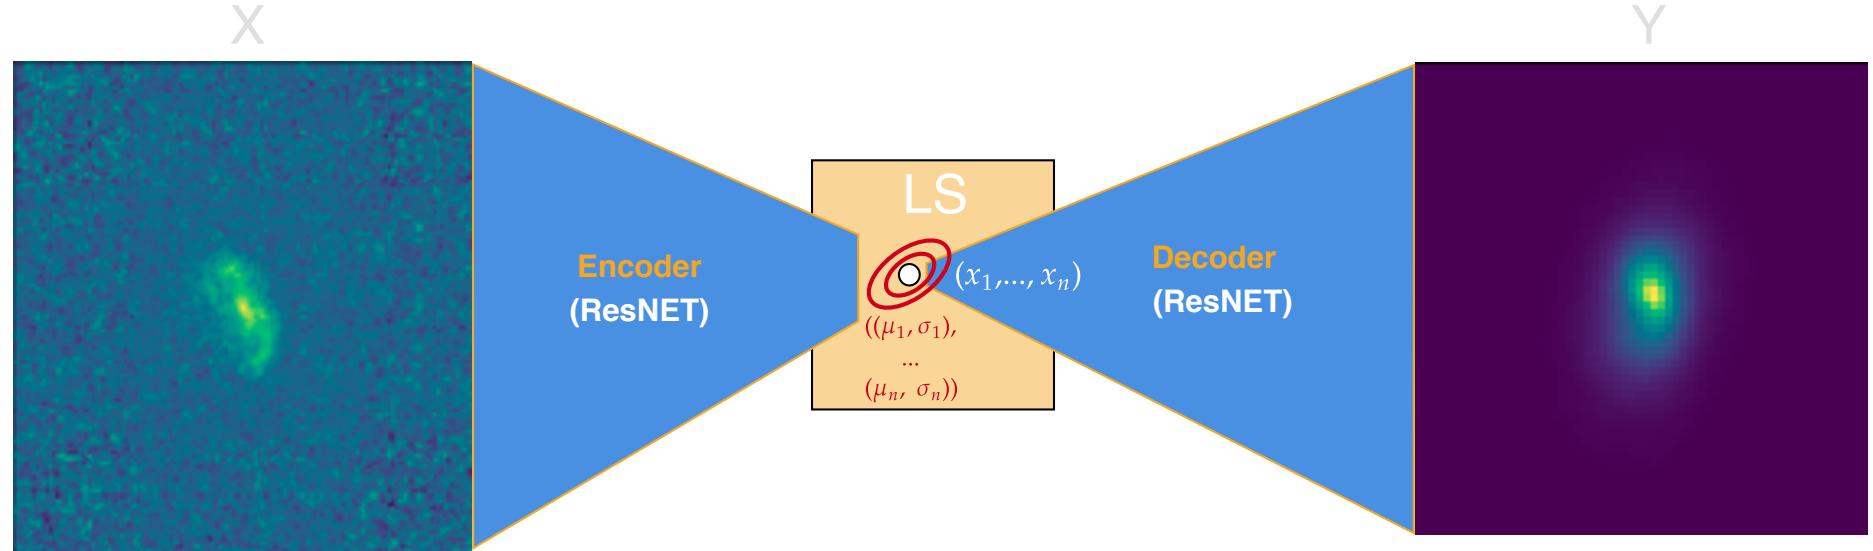
Others are more difficult to model

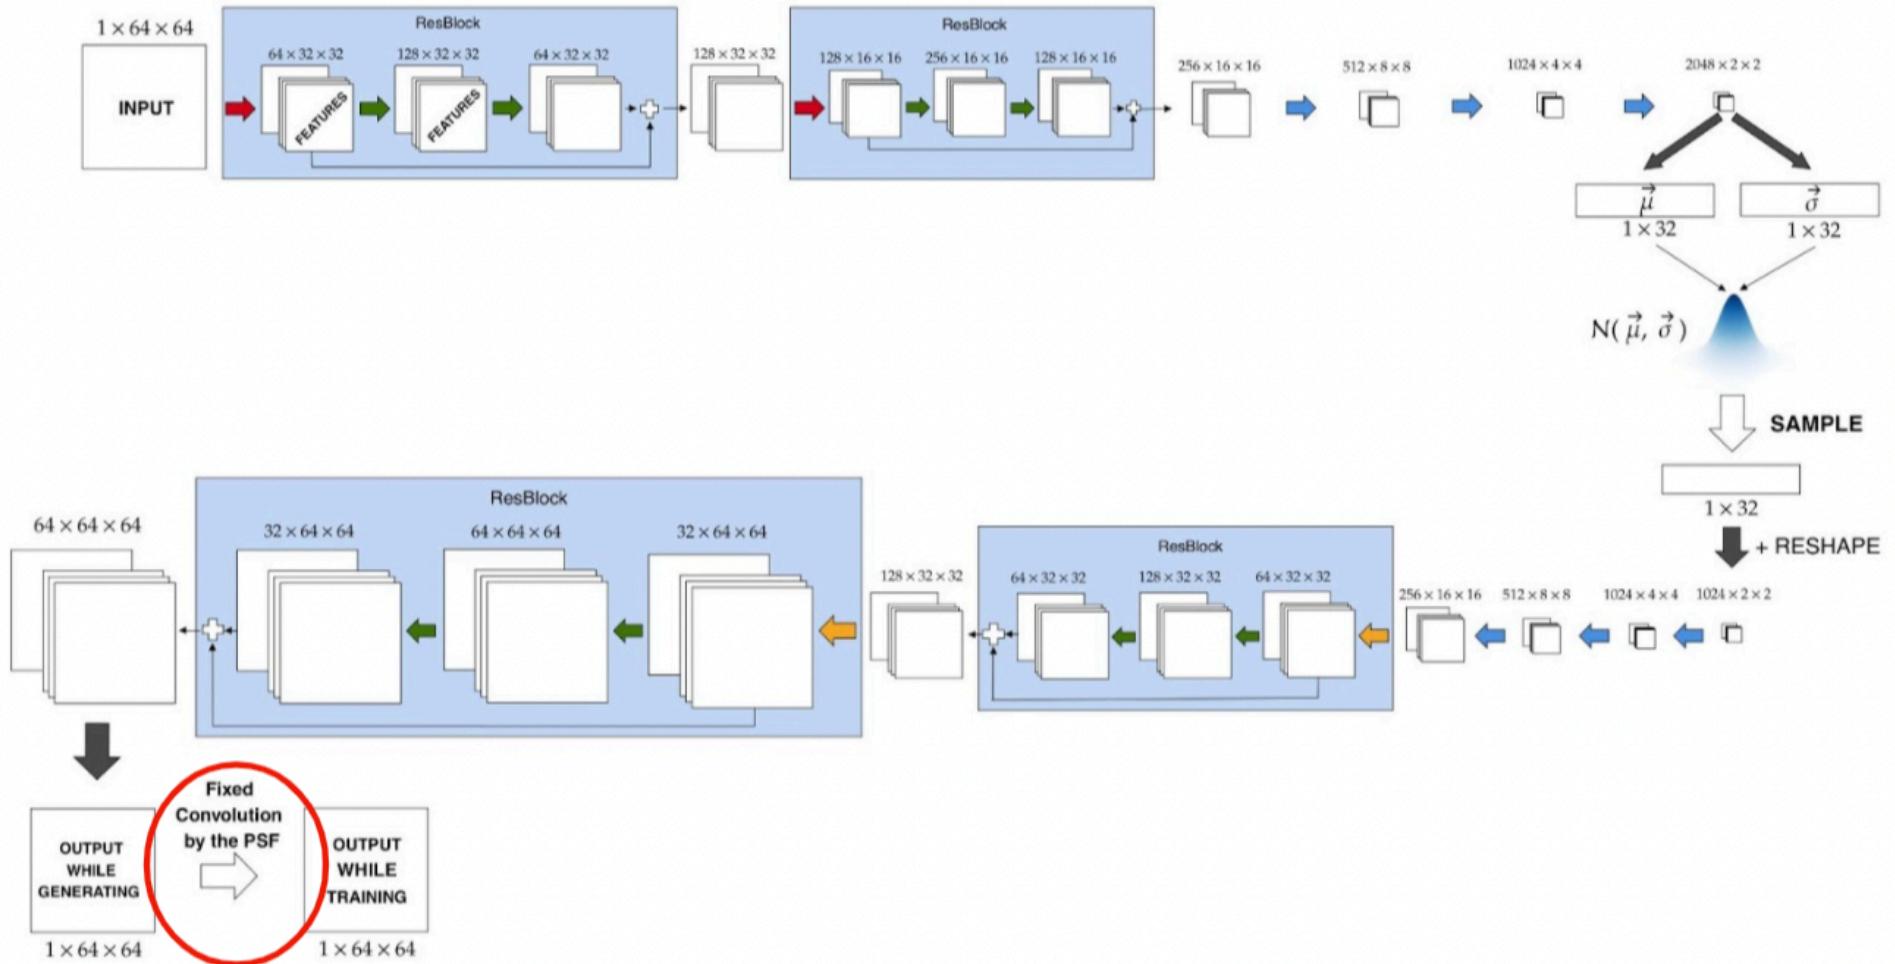
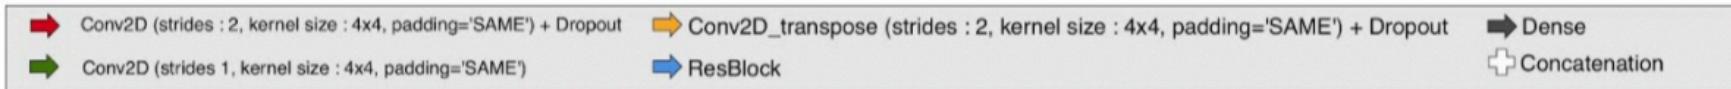


Include a Generative Model for galaxy generation in the Euclid Simulation Pipeline to model process for which we do not have a physical model



Include a Generative Model for galaxy generation in the Euclid Simulation Pipeline to model process for which we do not have a physical model

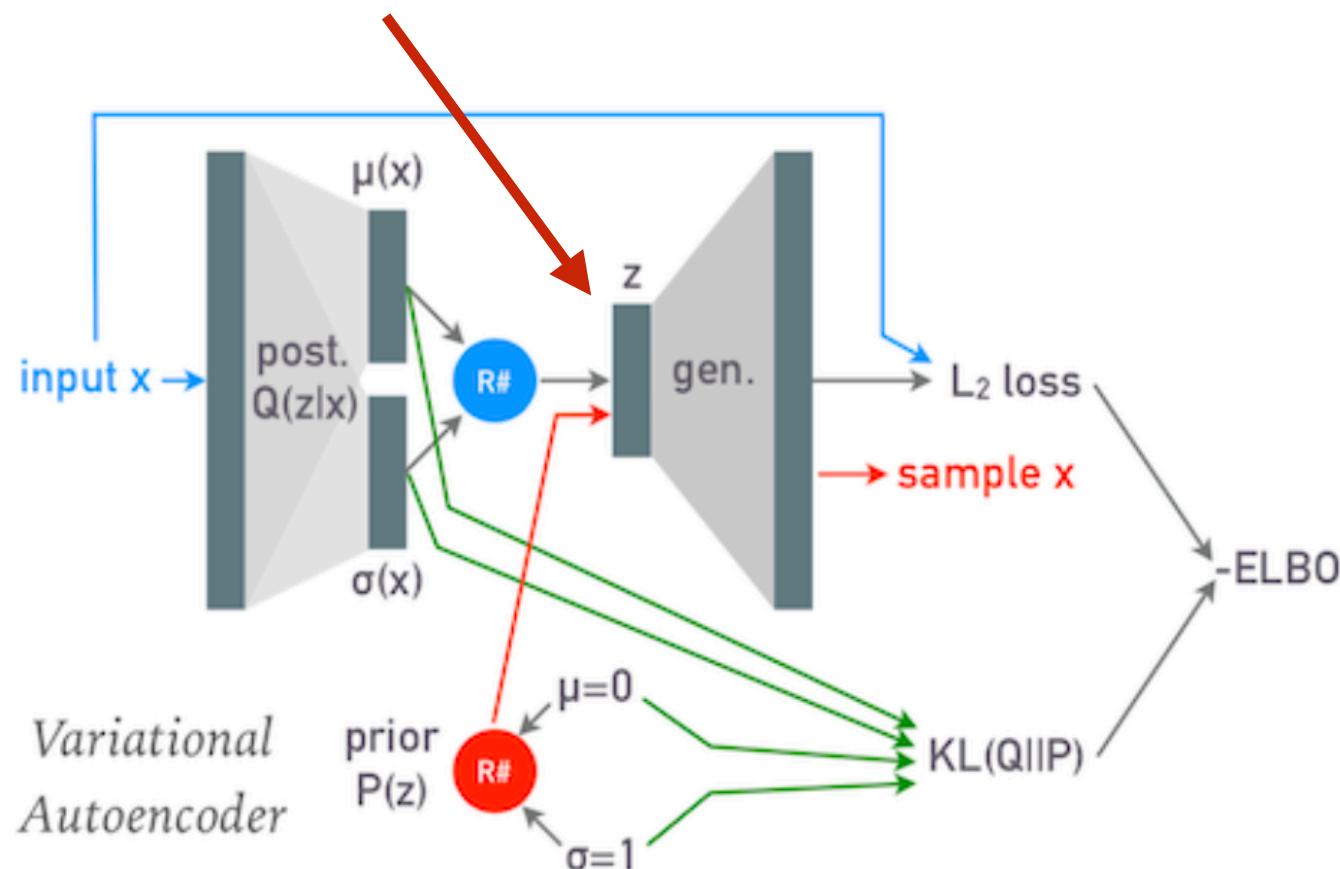


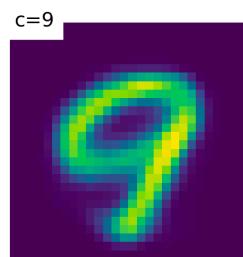
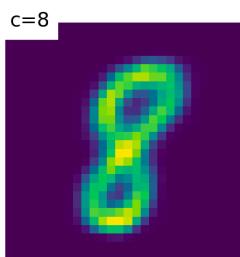
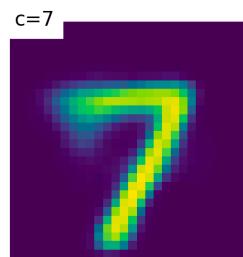
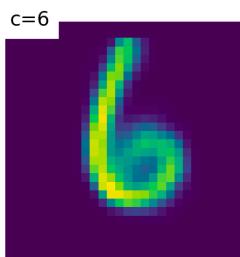
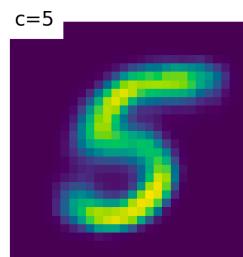
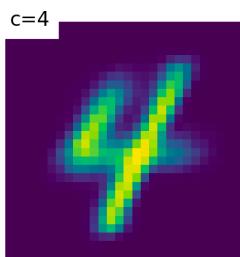
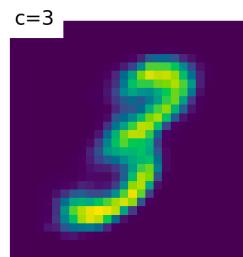
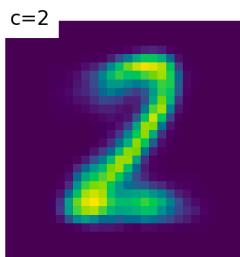
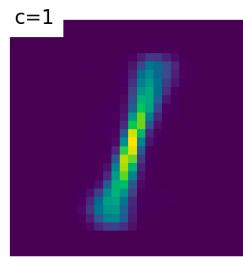
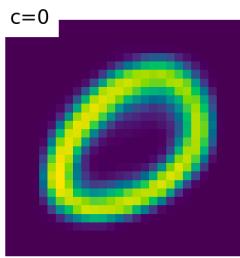


Physical information is included as well within the architecture

VAE's can also be conditioned, to generate a given class

concatenate labels





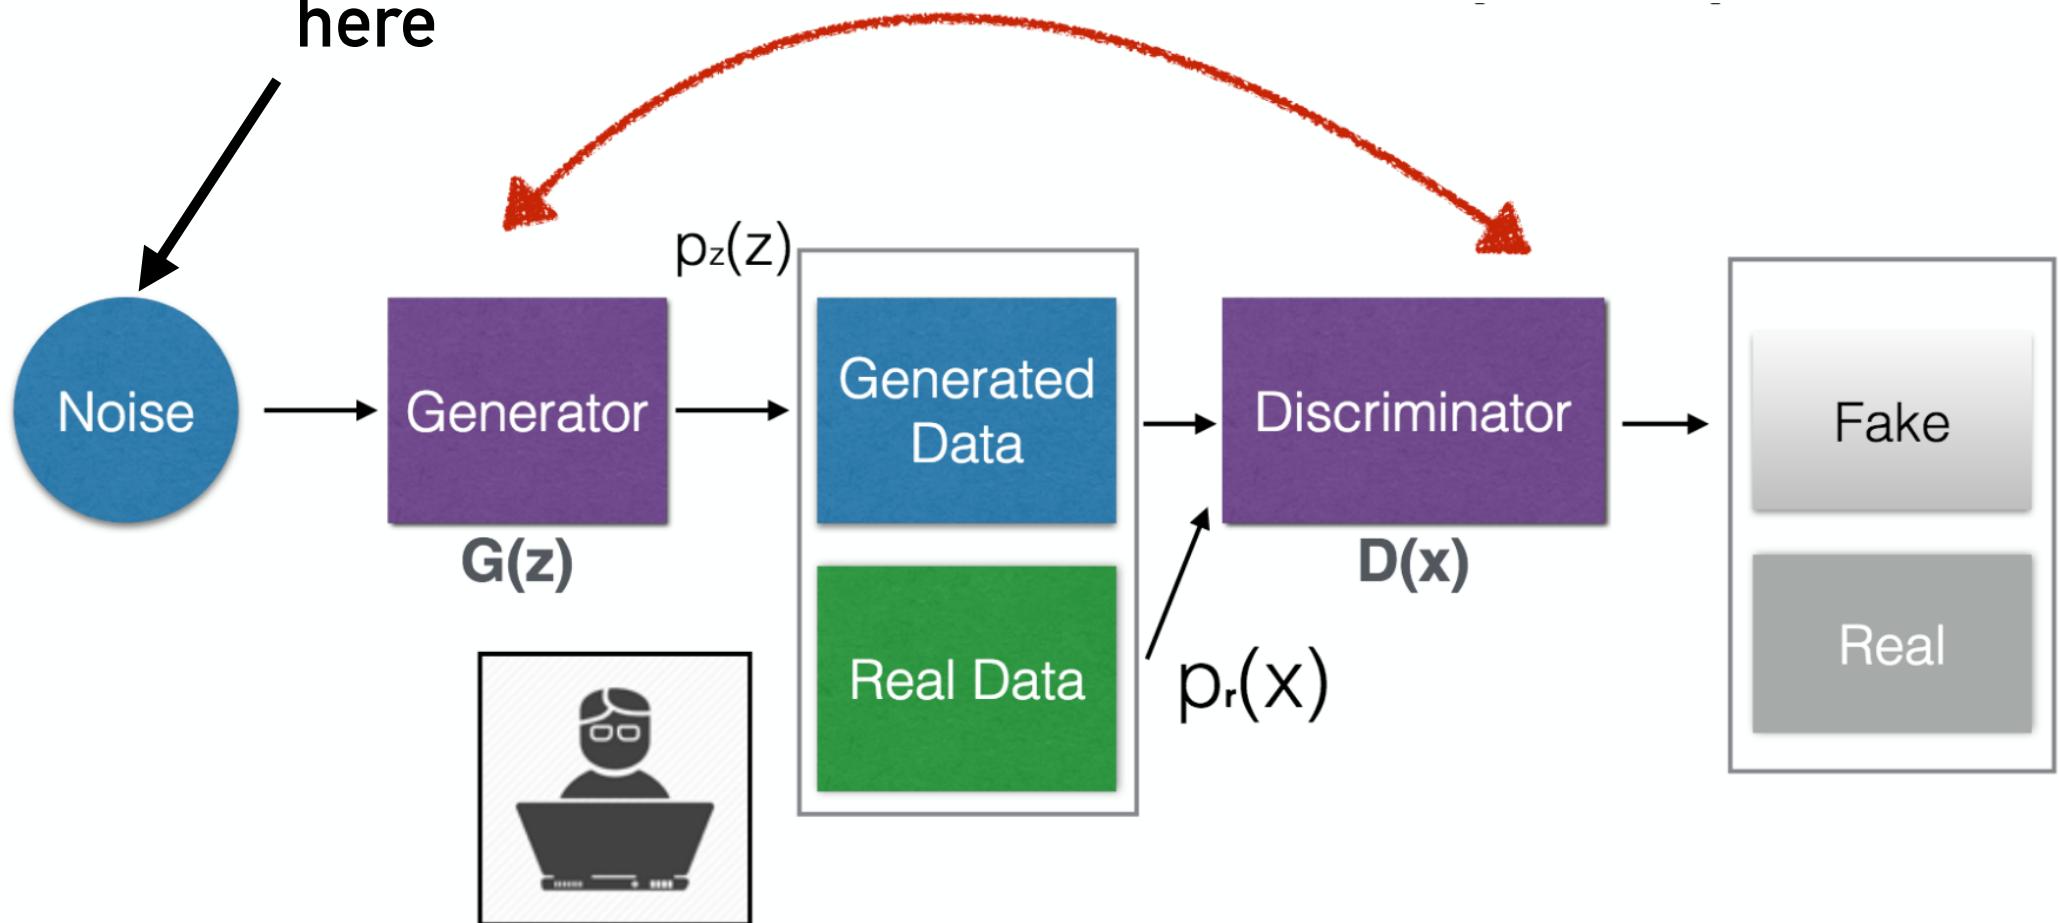
As for VAEs,
the goal of generative Adversarial Networks (GANs) is to estimate
 $p(z|x)$

They convert the problem into a “supervised approach” by using
two competing neural networks

GENERATIVE ADVERSARIAL NETWORKS

(Goodfellow+14)

The latent variable is
here

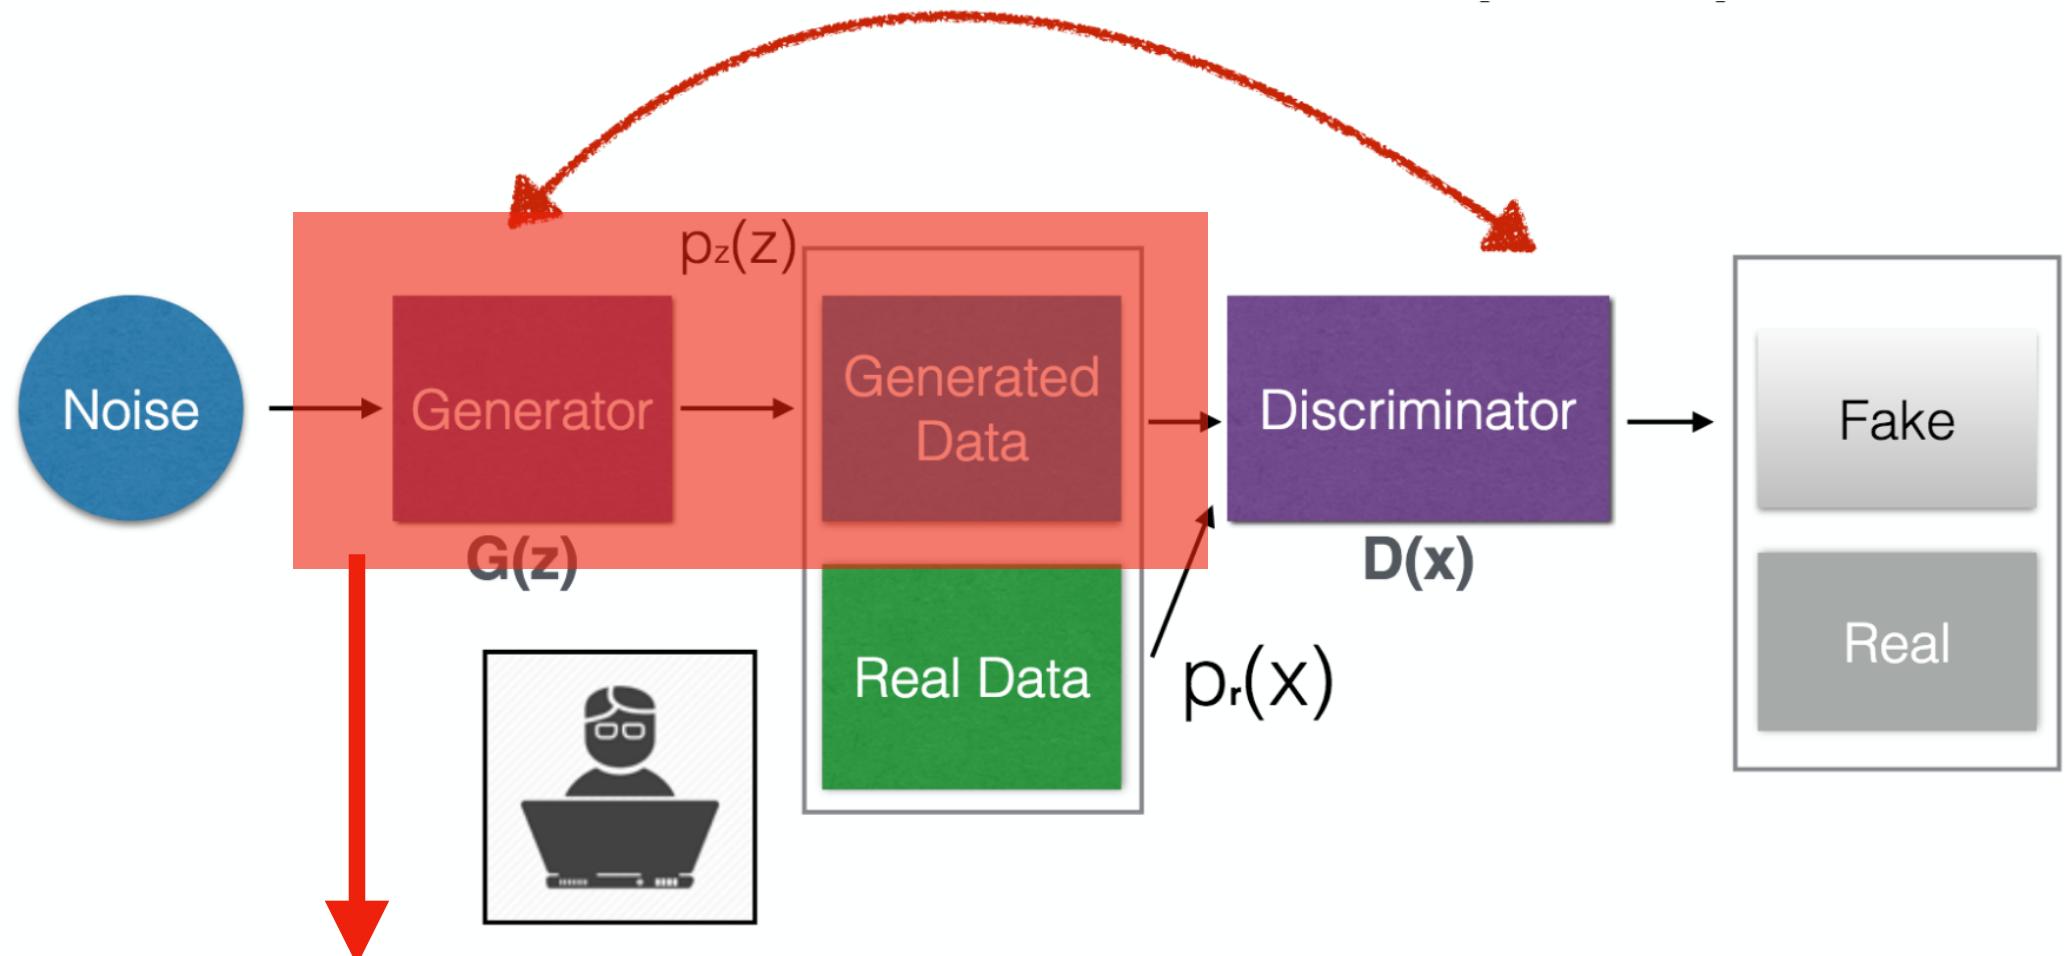


TWO COMPETING NETWORKS

GENERATIVE ADVERSARIAL NETWORKS

(Goodfellow+)

TWO COMPETING NETWORKS

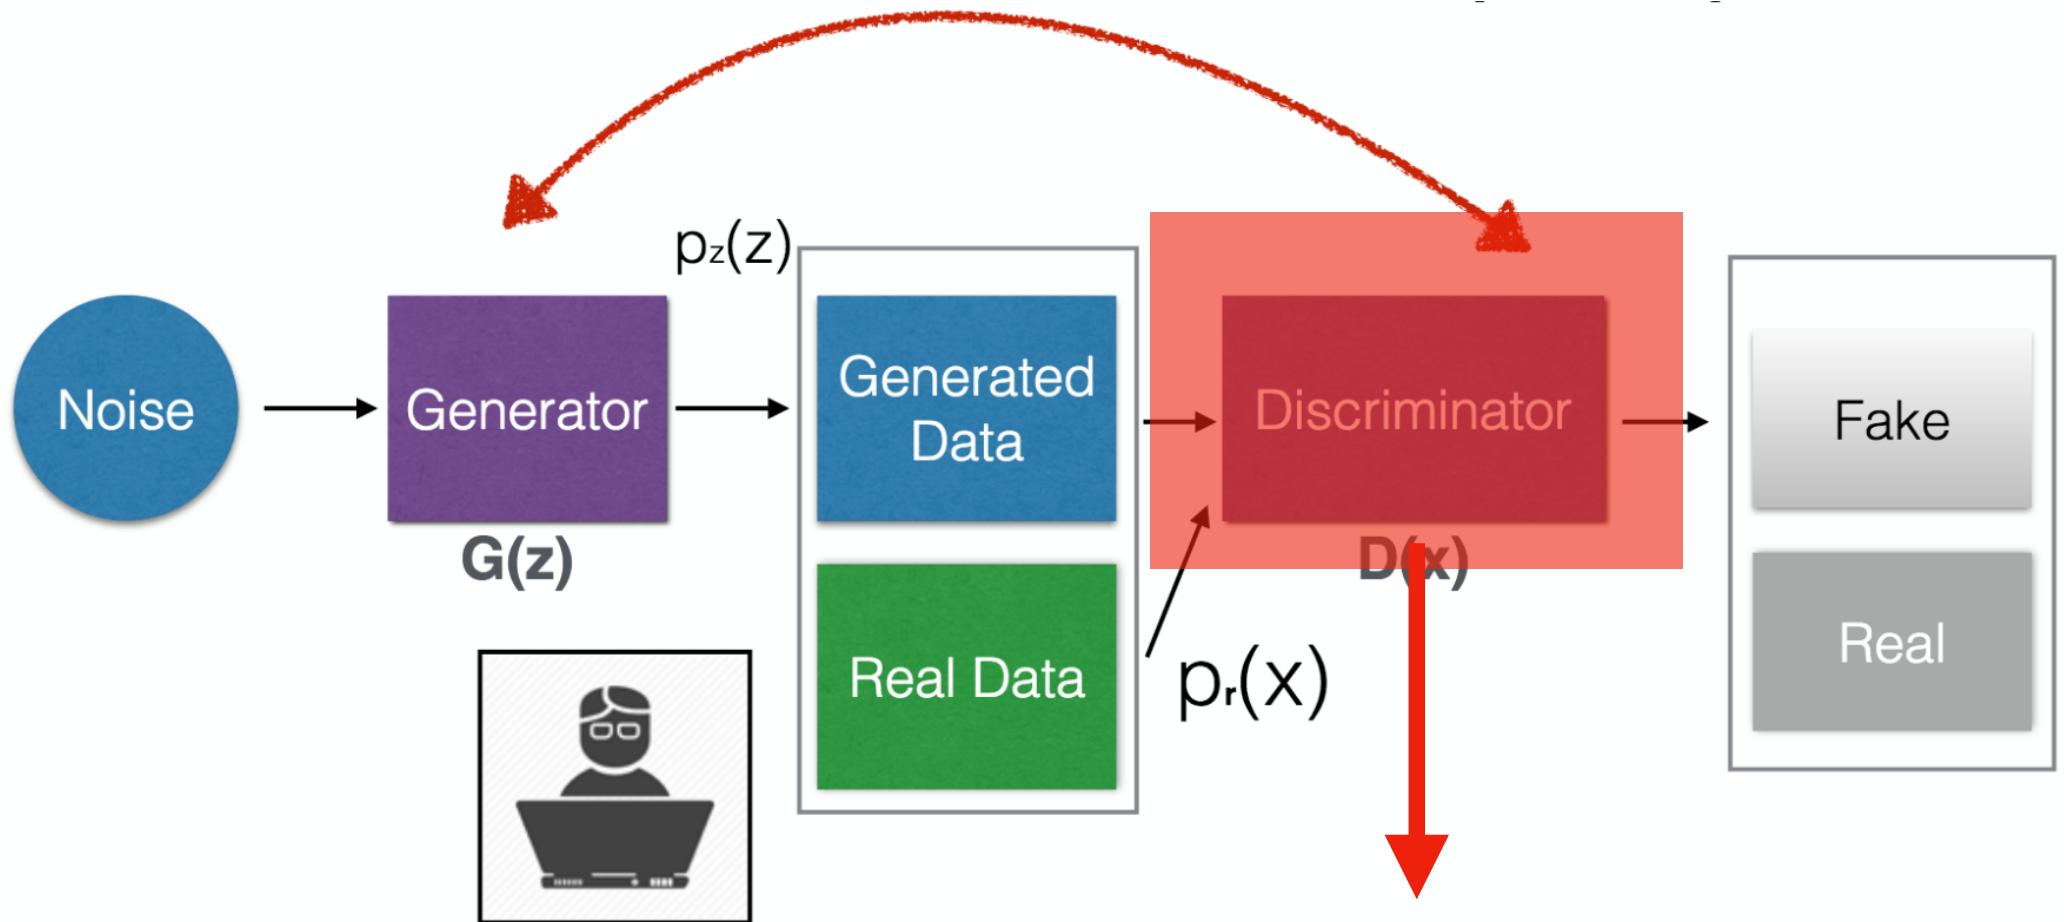


Every N iterations the generator
is trained to force the discriminator
to classify as real

GENERATIVE ADVERSARIAL NETWORKS

(Goodfellow+)

TWO COMPETING NETWORKS



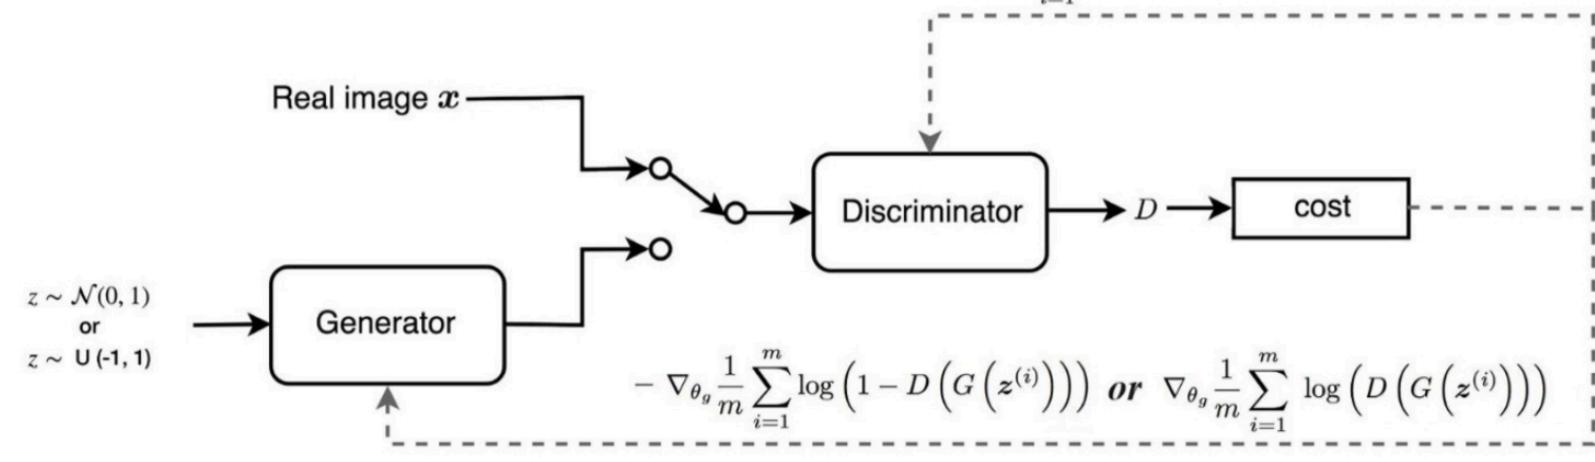
Every N iterations the discriminator
is trained to force to distinguish between
real and fake

IN PRACTICE

GAN LOSSES

DISCRIMINATOR LOSS (CROSS-ENTROPY)

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right]$$



GENERATOR LOSS

GANs CAN ACHIEVE IMPRESSIVE RESULTS...



Karras+19

THEY ARE ALSO VERY HARD TO TRAIN

1. IT IS HARD TO REACH EQUILIBRIUM AND IT IS ACTUALLY NOT GUARANTEED...

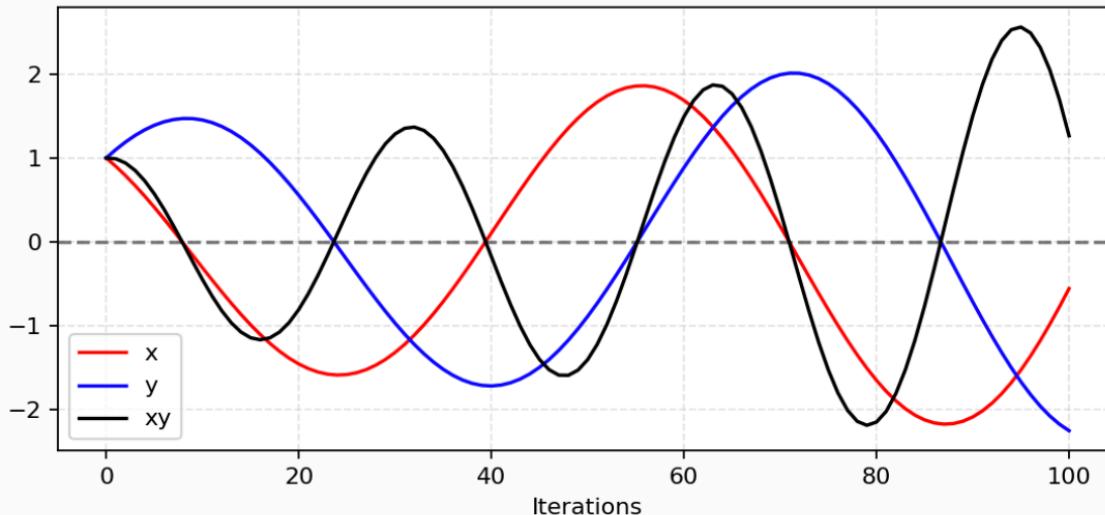
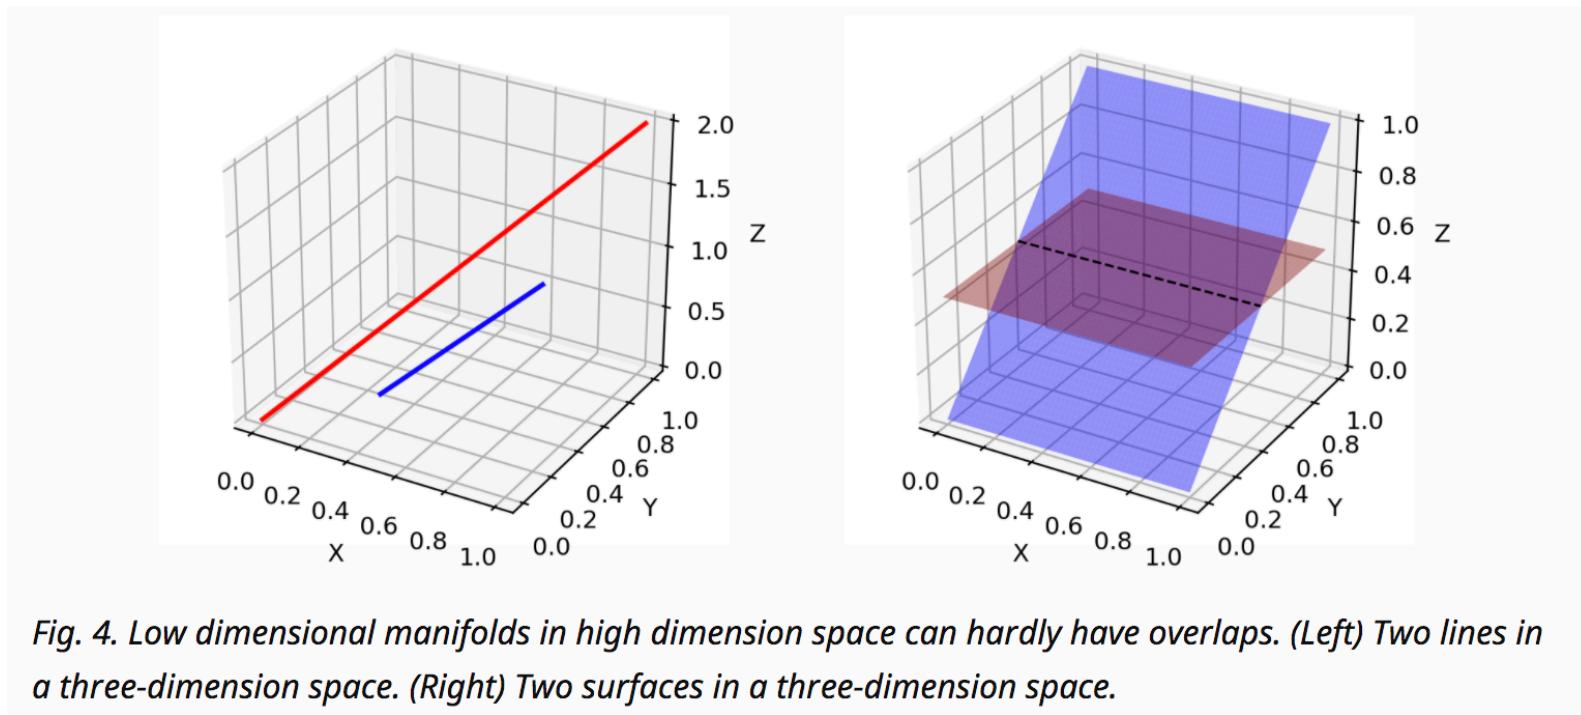


Fig. 3. A simulation of our example for updating x to minimize xy and updating y to minimize $-xy$. The learning rate $\eta = 0.1$. With more iterations, the oscillation grows more and more unstable.

THEY ARE ALSO VERY HARD TO TRAIN

2. LOW DIMENSIONAL SUPPORTS



THEY ARE ALSO VERY HARD TO TRAIN

3. MODE COLLAPSE

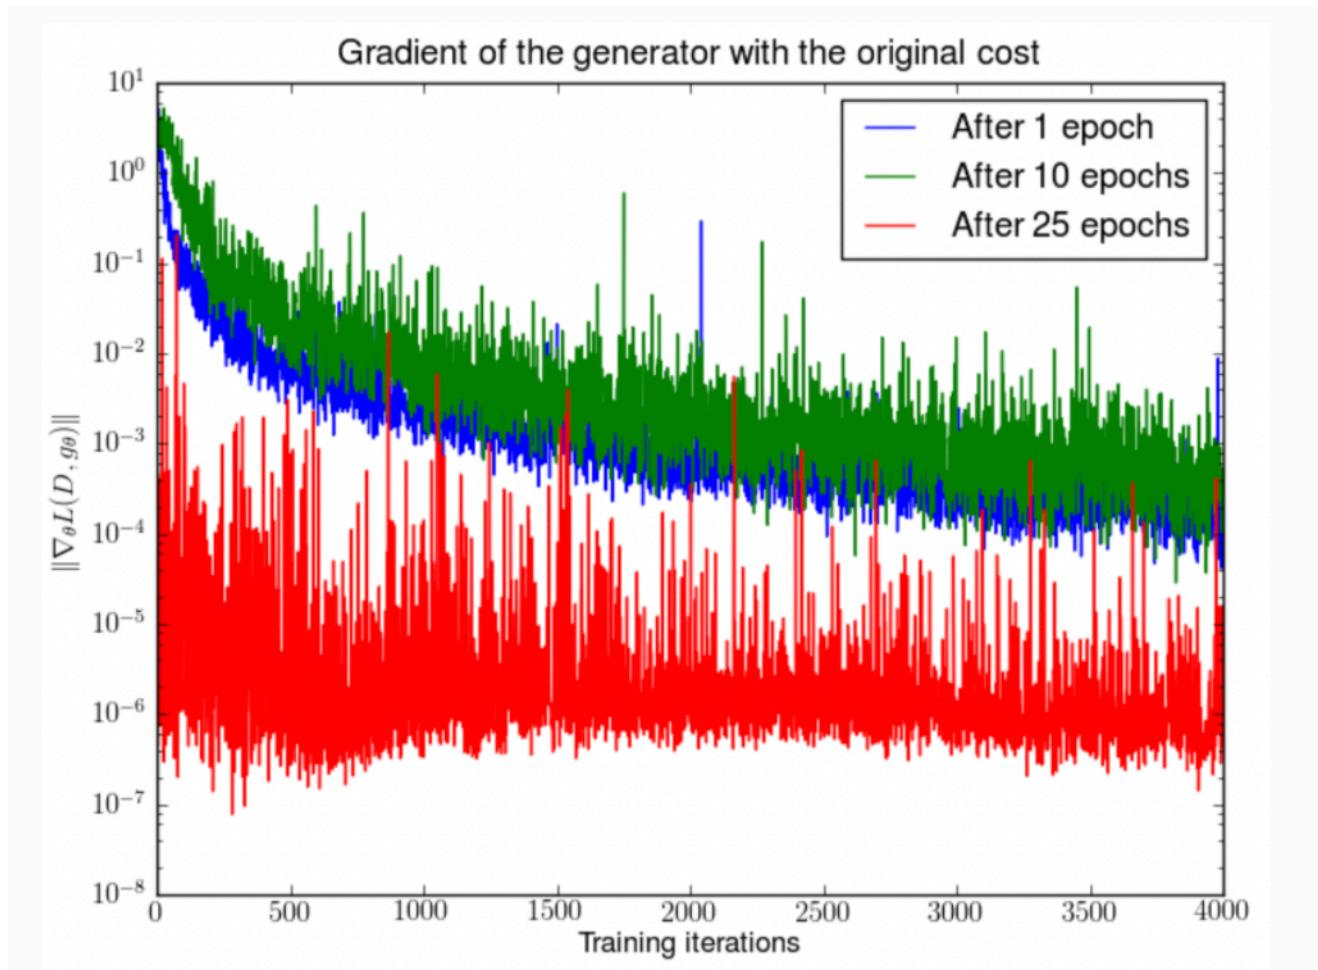


<https://lilianweng.github.io/lil-log/2017/08/20/from-GAN-to-WGAN.html>

razavi+19

THEY ARE ALSO VERY HARD TO TRAIN

4. VANISHING GRADIENT

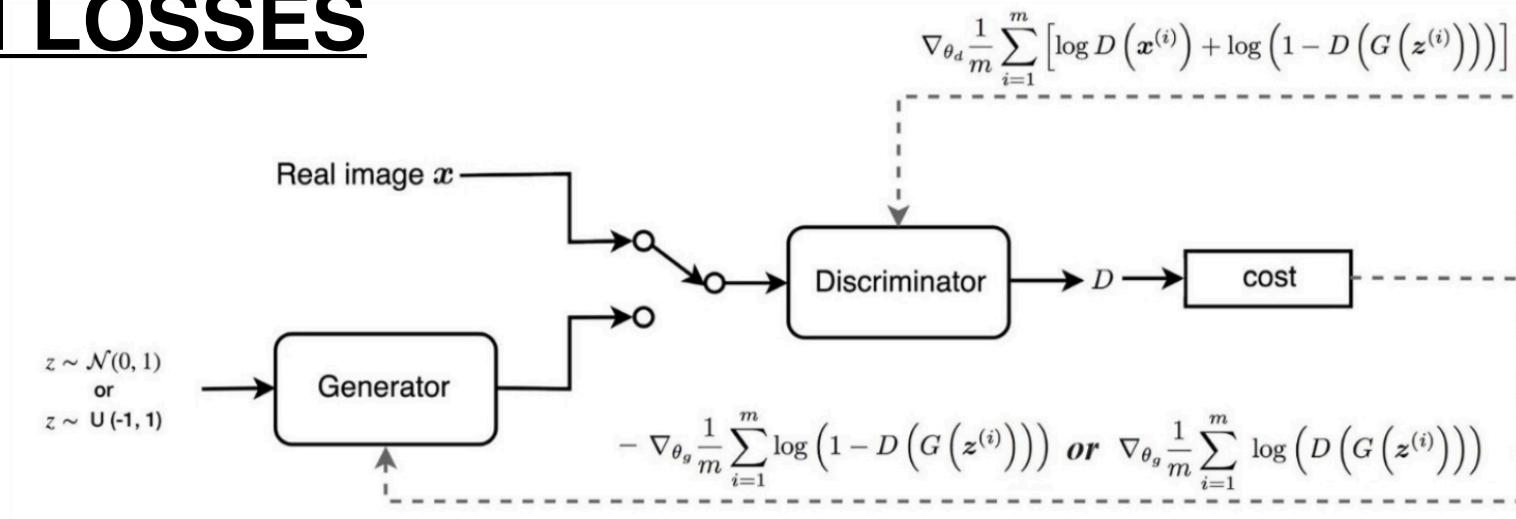


When the discriminator becomes perfect, there are no more gradients to update

IN PRACTICE

DISCRIMINATOR LOSS (CROSS-ENTROPY)

GAN LOSSES



GENERATOR LOSS

WASSERSTEIN GAN LOSSES

Cross entropy is replaced by Wasserstein distance

